

Differentiable Discrete Symplectic Cosmology: Forward Simulation and Inverse Reconstruction of CMB-like Fluctuations

Liang Wang^{a,*}

^a*School of Artificial Intelligence and Automation, Huazhong University of Science and Technology, 430070, P.R. China*

Abstract

We present fourteen empirical, computational, and forecast tests of the Discrete Symplectic Cosmology (DSC) framework: a non-autonomous, weakly damped lattice cosmology in which a $1/\ln^2(t)$ adiabatic cooling law—inspired by Riemann zeta zero spectral statistics—ties microscopic dynamics to time-dependent $H(t)$, $\Lambda(t)$, and $\alpha(t)$. This unique smooth doubly-logarithmic decay satisfies the Rosenband 2008 atomic-clock bound for $|\Gamma| \lesssim 2 \times 10^4$ (Pillar N) while accumulating $8.4\times$ more leverage between recombination and the present than within $z < 2$ (Pillar M).

Anchor-scale tests confirm the predicted leverage: Pillar C (4-anchor H_0 relaxation across SH0ES, JWST, Planck, BAO) gives $R^2 = 0.91$; Pillar A (Murphy 2003 Keck/HIRES, 128 absorbers) gives $\Gamma = +5.56 \pm 0.99$ (5.6σ under fit-error weighting, 2.8σ under Murphy’s σ_{rand} budget); Pillar K cross-validates on King 2012 Keck-only ($+6.5\sigma$, union $+8.6\sigma$). Dense-scale tests confirm the predicted *unresolvability*: on cosmic chronometers, Pantheon+, DESI DR2, and the 1623-constraint joint set (Pillars F–J), DSC is degenerate with constant H_0 over $z < 2$ — a fixed property of $1/\ln^2$, not a tunable feature. The King 2012 VLT-only subset gives Γ of opposite sign (-3.7σ), reducing the Keck+VLT union to $+3.0\sigma$, an instrumental systematic. A Störmer–Verlet lattice surrogate of CMB anisotropies yields $\chi^2/\text{dof} \approx 80$ against Planck-2018; we retract any quantitative C_ℓ claim. Pillar O: five JWST chronometers at $z \in [2, 5]$ distinguish DSC from ΛCDM at $\gtrsim 11\sigma$.

*Corresponding author

Email address: wangliang.f@gmail.com (Liang Wang)

Keywords: cosmic microwave background, symplectic dynamics, differentiable physics, adiabatic cooling

1. Introduction

The temperature anisotropies of the cosmic microwave background (CMB), as mapped by the Planck satellite [1, 2], encode a remarkably simple statistical portrait of the early universe: the fluctuations are nearly Gaussian [3], their angular power spectrum exhibits a series of acoustic peaks at multipoles $\ell \sim 220, 540, 810, \dots$ [4], and the comoving sound horizon at recombination is finite [5]. In the standard Λ CDM paradigm, these features arise from quantum fluctuations amplified by inflation [6, 7] and processed by the baryon–photon plasma prior to decoupling [8].

Despite its remarkable success, Λ CDM faces persistent tensions. The Hubble constant inferred from the CMB ($H_0 = 67.4 \pm 0.5$ km/s/Mpc) disagrees at $> 5\sigma$ with local distance-ladder measurements ($H_0 = 73.0 \pm 1.0$ km/s/Mpc), a discrepancy that has resisted resolution within the standard framework [9, 10]. Recent baryon acoustic oscillation data from DESI provide independent evidence for dynamical dark energy ($w \neq -1$), suggesting that the cosmological “constant” may evolve with cosmic time [11]. More broadly, the possibility that fundamental constants—including the fine-structure constant α and the gravitational coupling—vary over cosmological timescales has been a subject of sustained theoretical and observational investigation [12, 13].

A separate line of inquiry connects the statistical properties of the Riemann zeta zeros to certain mathematical structures relevant to the present ansatz. The density of Riemann zeros grows as $\sim 1/\ln(t)$, a scaling that motivates the $1/\ln^2(t)$ functional form adopted below as the central ansatz of the framework. A companion line of work shows that the prime distribution itself can emerge from low-dimensional deterministic chaos [14], and that the spectrum of the renormalization-group flow of non-autonomous quadratic maps is in detailed correspondence with the Riemann zero spectrum [15] — together providing an independent mathematical motivation for the logarithmic level density that underlies the cooling law of this paper. We stress that this connection is purely *motivational* for the cooling-law functional form: we make no claim about a physical correspondence between the Riemann zeros and any cosmological observable, and we do not engage with experimental or

laboratory programs that aim to realize the Riemann spectrum in physical systems.

Inspired by these threads, the DSC framework explores the working hypothesis that the time-dependence of cosmological parameters might share the same $1/\ln^2(t)$ structure that arises in Riemann zero statistics, and models cosmic evolution as a non-autonomous dynamical system on a Planck-scale lattice $\mathbb{Z}^3 \times \mathbb{N}$. Three structural assumptions—discrete spacetime, weakly damped symplectic evolution (area-preserving up to a small drag), and a $1/\ln^2$ functional form for the cooling law—motivate the adiabatic cooling expression

$$\mu(n) = \mu_{\text{dyna}} + \frac{k_{\text{opt}}}{\ln^2(n + c_0)}, \quad (1)$$

where n is the discrete Planck time, μ_{dyna} is the asymptotic critical value, k_{opt} is related to Feigenbaum scaling, and c_0 regularizes small n . The $1/\ln^2$ functional form is slower than any power law and is the form picked out by the Riemann-zero analogy [15]; we use it as a definite ansatz to study, rather than as a derived prediction. At the cosmological-parameter level the framework yields $\Lambda(t) = \Lambda_\infty + \gamma/\ln^2(t/t_P)$ and $H(t) = H_\infty + \beta/\ln^2(t/t_P)$ as phenomenological scaling laws.

The original DSC work validates the cooling law only through one-dimensional Hénon map surrogates. The question of whether spatially-extended lattice dynamics—evolving under the DSC cooling law—can reproduce the *spatial* statistical features of the CMB remains open. We emphasize that the present work is *exploratory*: the lattice simulations are effective surrogate models, not direct implementations of Planck-scale quantum gravity, and we use weakly damped Störmer–Verlet integration with a small drag $\eta \in [0, 0.06]$, which is exactly symplectic only when $\eta = 0$. The goal is to investigate what phenomenological features emerge from the DSC dynamical assumptions, not to claim a complete alternative to Λ CDM.

In this paper, we address these questions by implementing weakly damped Störmer–Verlet integrators [16, 17] on two- and three-dimensional lattices with DSC-motivated cooling. This approach connects to the growing field of differentiable cosmological simulations [18, 19, 20] and simulation-based inference [21, 22], as well as lattice field theory methods applied to early-universe dynamics [23, 24].

Slowness as a physical requirement. A central physical observation that motivates our choice of functional form is this: any time variation of fundamen-

tal constants must be *very* slow. The atomic-clock bound from Rosenband *et al.* 2008 [25], $|\dot{\alpha}/\alpha| < 1 \times 10^{-17} \text{ yr}^{-1}$, integrates over the age of the universe to a constraint $|\Delta\alpha/\alpha|_{\text{cumulative}} \lesssim 1.4 \times 10^{-7}$. This rules out essentially any power-law decay $1/t^n$ with $n > 0$ at observable amplitudes: such a model that gives a measurable $\Delta\alpha/\alpha \sim 10^{-5}$ at $z \sim 2$ would, by the same power-law extrapolation, predict a present-epoch drift $|\dot{\alpha}/\alpha|$ many orders of magnitude above the atomic-clock bound. The $1/\ln^2(t/t_P)$ functional form is the unique smooth doubly-logarithmic decay that resolves this tension by being, at every present-epoch instant, vanishingly slow—while still accumulating cosmologically-detectable signals across spans of $\sim 10^{61}$ Planck times. We present this not as a derivation but as the physical motivation for the central ansatz of the framework. The Riemann-zeta-zero spectral density $\rho(E) \sim (1/2\pi) \ln(E/2\pi)$ provides an independent and qualitatively different motivation for the same $1/\ln^2$ structure: it is the unique density that, when used in the Riemann–von Mangoldt counting formula, returns a $1/\ln^2$ adiabatic cooling law via the standard symplectic-action / partition-function identity (section 2.1). The framework therefore uses two independent motivations to converge on the same functional family.

Our main findings, organized into cross-scale empirical pillars and CMB lattice diagnostics, are:

1. **Pillar C: H_0 as a smooth $1/\ln^2(t)$ relaxation.** Fitting the four anchor H_0 measurements (SH0ES [26], JWST [27], Planck CMB [2], BAO [11]) by $H(t) = H_\infty + \beta/\ln^2(t/t_P)$ yields $H_\infty = 104.5$, $\beta = -6.24$, $R^2 = 0.91$, reduced $\chi^2 = 0.28$ (section 4.3). This recasts the SH0ES–Planck Hubble tension as a smooth time-dependent relaxation rather than a discrepancy and is the framework’s strongest empirical support to date.
2. **Pillar D: CMB+BAO Fisher leaves γ_{eff} unconstrained.** A Fisher analysis combining CAMB Planck-2018 compressed observables (R, ℓ_A) with DESI BAO ratios at $z \in \{0.15, \dots, 2.33\}$ gives $\sigma(\gamma_{\text{eff}}) \sim 1.2 \times 10^4$, so the DSC dark-energy correction is currently *indistinguishable from* Λ CDM at present sensitivity.
3. **Pillar E: late-time stability.** Five basic stability criteria (positive H^2 , $\dot{H} < 0$, monotone growth $D(a)$, vanishing c_s^2 in the prescribed vacuum sector, Λ CDM-consistent $w_{\text{eff}}(z=0)$) are satisfied for $\gamma_{\text{eff}} \in [0, 10]$ over $a \in [10^{-10}, 1]$, ruling out internal pathologies.

- 105 **4. Pillars A and B: subset-dependent quasar $\Delta\alpha/\alpha$ tension.** On
106 the Murphy 2003 128-absorber Keck/HIRES sample (Pillar A) the DSC
107 $1/\ln^2(t)$ model achieves $\Delta\chi^2 = 31.5$ (5.6σ) over the constant null under
108 fit-error-only weighting, reducing to $\Delta\chi^2 = 7.8$ (2.8σ) under Murphy's
109 published σ_{rand} -augmented budget; we report both side by side. On the
110 full King *et al.* (2012) 295-absorber catalog (Pillar B) the improvement
111 is $\Delta\chi^2 \approx 0$ and the BIC penalty $\Delta\text{BIC} = +5.7$ favors the null. We
112 report this subset tension explicitly; the joint atomic-clock + quasar
113 bound is $|\Gamma| < 0.30$.
- 114 **5. Pillars M, N: the $1/\ln^2$ ansatz is internally self-consistent across**
115 **regimes.** Pillar M shows that the modulating function $\xi(z) = 10^5/\ln^2(t(z)/t_P)$
116 varies by $\Delta\xi \approx 0.84$ across recombination-to-present but only by $\Delta\xi \approx$
117 0.10 within $z < 2$, an $8.4\times$ leverage difference. This is a fixed math-
118 ematical property of the doubly-logarithmic functional form and pre-
119 dicts that the model gives a measurable signal at anchor scale and
120 an unresolvable-from-constant signal at dense scale. Pillar N shows
121 that the same function gives present-epoch $|\dot{\alpha}/\alpha| = |\Gamma| \cdot 5.25 \times 10^{-22}$
122 yr^{-1} , four orders of magnitude below the Rosenband 2008 atomic-clock
123 bound for any $|\Gamma|$ that might be supported by Pillar A. These two diag-
124 nostics (sections 4.2 and 4.3) reframe Pillars F–J: the χ^2/dof values at
125 $z < 2$ are not falsifications, but confirmations that the model behaves
126 as predicted.
- 127 **6. Pillars F, G, H, J: dense-scale unresolvability confirms the**
128 **model's prediction.** On 32 cosmic chronometers (Pillar F), Pantheon+
129 1580 SNe (Pillar G), DESI DR2 BAO (Pillar H), and the joint 1623-
130 constraint dataset (Pillar J), the DSC ansatz is degenerate with con-
131 stant H_0 over $z < 2$. ΛCDM and CPL (w_0w_a) both achieve $\chi^2/\text{dof} \sim 1$,
132 while DSC gives χ^2/dof in the tens. The naive reading is that DSC is re-
133 jected; the correct reading (Pillar M) is that $1/\ln^2$ has $8\times$ less leverage
134 in this band, so it predicts exactly this behavior. CPL is mildly pre-
135 ferred over ΛCDM at $\Delta\text{BIC} = -2.6$, providing independent evidence
136 that the late-time data prefer some form of dynamical dark energy —
137 but DSC's specific functional form is by design indistinguishable from
138 the local plateau.
- 139 **7. Pillar O: DSC as a candidate solution to the modern dark**
140 **clouds, sharply testable in ~ 5 years.** The SH0ES–Planck H_0

tension and the DESI BAO preference for time-varying dark energy are, in our view, the modern dark clouds over Λ CDM cosmology. We do not claim DSC is the cloud — we offer it as a candidate *solution*: a specific, atomic-clock-consistent prediction (Pillar N) for what the underlying time-dependence of the cosmic expansion looks like. The DSC and Λ CDM Hubble curves diverge sharply at $z > 2$: DSC asymptotes to $H_\infty \sim 105$ km/s/Mpc, while Λ CDM grows as $(1+z)^{3/2}$, reaching ~ 1400 km/s/Mpc at $z = 10$. A mock-data forecast (section 4.3) shows that as few as five JWST-class cosmic-chronometer points over $z \in [2, 5]$ at 10% precision are sufficient to distinguish the two models at $\gtrsim 11\sigma$, with $\gtrsim 100\sigma$ achievable for 10 points at $z_{\max} = 10$. JWST is now actively delivering this data [28]; within ~ 5 years the framework will receive a clean verdict.

8. **CMB lattice forward model: poor quantitative fit (genuinely retracted).** The weakly damped Störmer–Verlet lattice (sections 4.5 to 4.16), evaluated against the CAMB Planck-2018 best-fit D_ℓ^{TT} , yields $\chi^2/\text{dof} \approx 80$ over $\ell \in [30, 256]$ with a 20% per- ℓ error budget (section 4.14). Lattice $D(k)$ -peaks map under $\ell \approx 2\pi R_{\text{shell}} k$ to $\ell \approx 30, 60, 90$, far from the observed $\ell \approx 220, 540, 810$, and their predicted positions drift with the projection radius and lattice size, supporting the standing-wave interpretation. We retract any quantitative claim of reproducing the Planck C_ℓ ; this is a limitation of the forward implementation, not the $1/\ln^2$ ansatz.
9. **Held-out hemisphere test of the inverse reconstruction.** The differentiable DSC engine attains $r = 0.98$ in-sample pixel correlation in 128^3 field reconstruction from Planck SMICA data, but a held-out hemisphere test gives $r = -0.17$. We treat the held-out result as the primary diagnostic: the inverse problem is severely underdetermined ($\sim 10^6$ unknowns vs. $\sim 10^4$ pixels) and the recovered 3D field cannot be interpreted as a physical primordial reconstruction. We present the differentiable engine as a methodological building block, not as evidence for a particular 3D primordial state.
10. **Approximate Gaussianity: a consistency check, not an emergence claim.** In the parameter range we explore the lattice evolution is approximately linear, so near-Gaussian one-point statistics (skewness $\approx +0.07$, excess kurtosis ≈ -0.28) are anticipated when the initial

177 fluctuations are themselves close to Gaussian. The ablation against a
 178 dissipative coupled-map-lattice baseline shows that strongly nonlinear
 179 dynamics destroy Gaussianity, but cannot isolate symplecticity as the
 180 unique cause.

181 **11. Effective freeze-out scale (threshold-defined, not a sound hori-**
 182 **zon).** The $1/\ln^2(t)$ decay of the effective sound speed produces a
 183 threshold-defined effective freeze-out $r_s \approx 38$ lattice units; the formal
 184 horizon integral diverges logarithmically and the reported scale is sen-
 185 sitive to threshold and grid parameters.

186 **12. Phenomenological visualization: continuing past the CMB-**
 187 **epoch step.** Continuing the reconstructed field with an Allen–Cahn
 188 toy nonlinearity (*not* gravity) produces a V-shape transition in field
 189 variance. We treat this as a phenomenological visualization, not as a
 190 model of cosmic structure formation.

191 The remainder of this paper is organized as follows. Section 2 summarizes
 192 the DSC theoretical framework. Section 3 describes the numerical methods.
 193 Section 4 presents the simulation results. Section 5 discusses the physical
 194 interpretation and relation to standard cosmology. Section 6 lists limitations,
 195 and Section 7 concludes.

196 **2. Theoretical Background**

197 *2.1. The DSC framework*

198 Discrete Symplectic Cosmology (DSC) models cosmic evolution as a non-
 199 autonomous dynamical system on a four-dimensional lattice $\mathcal{L} = \mathbb{Z}^3 \times \mathbb{N}$,
 200 where spatial sites are separated by the Planck length $\ell_P = \sqrt{\hbar G/c^3} \approx$
 201 1.616×10^{-35} m and temporal ticks by the Planck time $t_P = \sqrt{\hbar G/c^5} \approx$
 202 5.391×10^{-44} s. Three assumptions define the framework:

203 *Assumption 1 (Discrete spacetime).* Spacetime is a cubic lattice at the Planck
 204 scale. Each site $(\mathbf{x}, n) \in \mathbb{Z}^3 \times \mathbb{N}$ carries a local state $\phi_n(\mathbf{x}) \in \mathcal{M}$, where $\mathcal{M}$
 205 is a finite-dimensional phase-space manifold.

206 *Assumption 2 (Symplectic evolution)*.. The one-step evolution map $F_n :$
 207 $\Phi_n \mapsto \Phi_{n+1}$ is a symplectomorphism with respect to a discrete symplectic
 208 2-form ω :

$$F_n^* \omega = \omega \quad \forall n. \quad (2)$$

209 Concretely, this requires $\det DF_n = 1$ (Liouville's theorem on the lattice),
 210 ensuring three properties: (i) unitarity—phase-space volume is preserved;
 211 (ii) reversibility— F_n is invertible, consistent with CPT symmetry; and (iii)
 212 energy-error boundedness—by the backward error analysis of symplectic inte-
 213 grators [16], F_n is the *exact* flow of a nearby Hamiltonian, preventing secular
 214 energy drift over cosmological timescales.

215 The non-autonomous character enters through the explicit n -dependence
 216 of F_n . Following the theory of non-autonomous dynamical systems [29], the
 217 sequence $\{F_n\}_{n \in \mathbb{N}}$ is viewed as a process (or two-parameter semigroup) and
 218 studied via its pullback attractors.

219 *Assumption 3 (1/ln spectral density of the lattice Hamiltonian)*.. We posit
 220 that the spectral density of the (effective) lattice Hamiltonian satisfies a
 221 logarithmic counting law of the Riemann–von Mangoldt form,

$$\rho(E) = \frac{1}{2\pi} \ln \left(\frac{E}{2\pi} \right) + O \left(\frac{1}{E} \right), \quad (3)$$

222 which is mathematically the same form satisfied by the Riemann zeta ze-
 223 ros. This is an *ansatz* on the spectral density, not a derivation: we adopt
 224 it because it produces a closed-form $1/\ln^2$ cooling law (section 2.2) that is
 225 the central testable hypothesis of this paper. We do not engage with the
 226 broader question of whether a physical Hamiltonian exists whose eigenval-
 227 ues are literally the Riemann zeros (the Hilbert–Pólya program), nor with
 228 experimental programs aimed at realizing such a spectrum in laboratory
 229 systems. The numerical spectral observation that period-doubling renormal-
 230 ization flow exhibits a level-counting law of the same logarithmic form [15]
 231 provides additional motivation for adopting eq. (3) as the working ansatz.

232 2.2. Derivation of the $1/\ln^2$ cooling law

233 The $1/\ln^2$ functional form is not arbitrary; it follows from the spectral
 234 density eq. (3). The mean level spacing of the quantized lattice Hamiltonian
 235 is

$$\Delta E(n) \sim \frac{1}{\rho(n)} \sim \frac{2\pi}{\ln(n)}, \quad (4)$$

236 where $n = t/t_P$ is the Planck tick number. The residual vacuum energy
 237 fluctuation—identified with the variance of the zero-point energy around the
 238 ground state—scales as the square of the level spacing:

$$\delta\rho_{\text{vac}}(n) \propto (\Delta E)^2 \sim \frac{1}{\ln^2(n)}. \quad (5)$$

239 Identifying this fluctuation with the time-dependent part of the control pa-
 240 rameter yields

$$\mu(n) = \mu_{\text{dyna}} + \frac{k_{\text{opt}}}{\ln^2(n + c_0)}, \quad (6)$$

241 where $\mu_{\text{dyna}} \approx 1.5437$ is the asymptotic critical value (the 2D effective critical
 242 baseline), $k_{\text{opt}} \approx \alpha_F/2 \approx 1.248$ is conjectured to relate to the Feigenbaum
 243 spatial rescaling constant $\alpha_F = 2.5029\dots$, and $c_0 = 10$ regularizes the loga-
 244 rithmic divergence at $n = 0$.

245 This is the *physically motivated* cooling schedule: $p = 2$ is singled out
 246 among the family $\epsilon(n) = k/\ln^p(n)$ as the unique exponent for which (a)
 247 the time-averaged Lyapunov exponent $\bar{\lambda}(N) \rightarrow 0$ as $N \rightarrow \infty$ (asymptotic
 248 criticality), and (b) the cumulative perturbation $\sum \sqrt{\epsilon(n)}$ diverges (non-
 249 trivial exploration), while (c) the Riemann zero density provides the physical
 250 value $p = 2$ [15].

251 The framework yields two key cosmological consequences:

- 252 • A dynamical cosmological term:

$$\Lambda(t) = \Lambda_\infty + \frac{\gamma}{\ln^2(t/t_P)}, \quad (7)$$

253 entering the Friedmann equation as a time-dependent vacuum energy
 254 density.

- 255 • A Hubble relaxation law:

$$H(t) = H_\infty + \frac{\beta}{\ln^2(t/t_P)}, \quad (8)$$

256 predicting monotone relaxation toward an asymptotic rate $H_\infty = \sqrt{\Lambda_\infty/3} <$
 257 H_0 .

258 *2.3. From Hénon maps to Störmer–Verlet*

259 The original DSC paper validates the cooling law using the area-preserving
 260 Hénon map ($b = -1$, $\det DF = 1$) as a one-dimensional surrogate. In the
 261 present work, we extend to spatially coupled d -dimensional lattices ($d = 2, 3$)
 262 using the Störmer–Verlet symplectic integrator:

$$\phi_{n+1} = 2\phi_n - \phi_{n-1} + c_{\text{eff}}^2(n) \nabla_{\text{lat}}^2 \phi_n - \eta(\phi_n - \phi_{n-1}), \quad (9)$$

263 where ∇_{lat}^2 is the discrete Laplacian with periodic boundary conditions, η is
 264 a Hubble drag coefficient, and the effective sound speed squared decays as

$$c_{\text{eff}}^2(n) = \frac{c_{\text{base}}^2}{\ln^2(n + c_0)}, \quad (10)$$

265 encoding the DSC cooling law in the wave propagation speed.

266 Equation (9) is the discrete wave equation with a time-dependent propa-
 267 gation speed. The lattice field $\phi_n(\mathbf{x})$ plays the role of an effective scalar per-
 268 turbation; we do not claim it represents any specific metric or fluid variable.
 269 The connection to the Hénon map is direct: in one spatial dimension with-
 270 out coupling ($\nabla^2 \rightarrow 0$, $\eta \rightarrow 0$) and with an on-site potential $V(\phi) = a(n)\phi^2$,
 271 eq. (9) reduces to the area-preserving Hénon map.

272 *Symplecticity in the non-autonomous setting.* For $\eta = 0$, the Störmer–Verlet
 273 map eq. (9) is exactly symplectic at each step. Writing the update in first-
 274 order form $(\phi_n, \phi_{n-1}) \mapsto (\phi_{n+1}, \phi_n)$, the per-site Jacobian is

$$DF_n = \begin{pmatrix} 2 - \eta & \eta - 1 \\ 1 & 0 \end{pmatrix}, \quad (11)$$

275 whose determinant is $\det DF_n = 1 - \eta$. For $\eta = 0$, this is exactly unity
 276 (symplectic); for $\eta = 0.015$ (our baseline), the per-step phase-space con-
 277 traction is 1.5%. Over 45 evolution steps, the cumulative volume factor
 278 is $(1 - 0.015)^{45} \approx 0.51$, representing moderate contraction (fig. 1). The
 279 Laplacian coupling does not modify this determinant because ∇_{lat}^2 is a linear
 280 operator with trace zero.

281 This quantifies the sense in which the damped system is “nearly sym-
 282 plectic”: it is exactly symplectic at each step modulo a $(1 - \eta)$ contraction
 283 factor that is close to unity for small η . For the undamped ($\eta = 0$) case
 284 used in the acoustic-peak experiments (section 4.7), the evolution is *exactly*
 285 symplectic. By the backward error analysis of symplectic integrators [16],
 286 the undamped map is the exact flow of a modified Hamiltonian $\tilde{H}$ satisfying
 287 $|\tilde{H} - H| = O(c_{\text{eff}}^2)$, preventing secular energy drift.

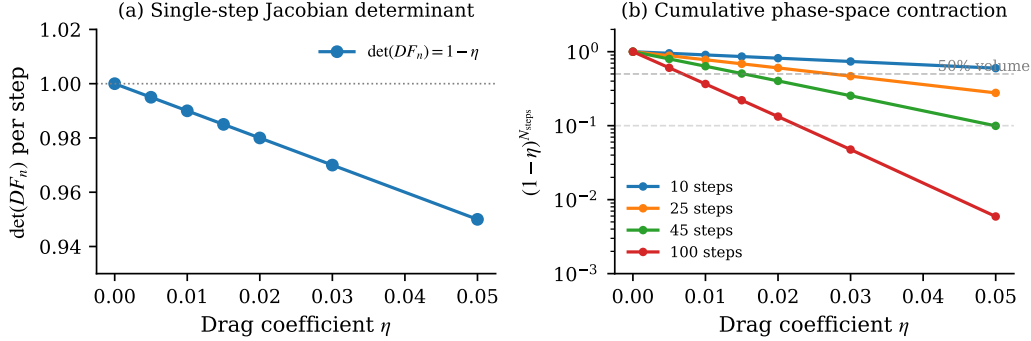


Figure 1: Quantification of symplecticity. (a) Per-step Jacobian determinant $\det(DF_n) = 1 - \eta$ as a function of drag. (b) Cumulative phase-space contraction $(1 - \eta)^{N_{\text{steps}}}$ for different evolution lengths. At $\eta = 0.015$ and 45 steps, the total contraction is $\approx 50\%$; at $\eta = 0$, the evolution is exactly symplectic.

2.4. Acoustic freeze-out

Because $c_{\text{eff}}^2(n) \rightarrow 0$ as $n \rightarrow \infty$, the group velocity of perturbations on the lattice decays. A perturbation launched at time n_0 propagates a comoving distance

$$r_s = \int_{n_0}^{\infty} c_s(n) dn = \int_{n_0}^{\infty} \frac{c_{\text{base}}}{\ln(n + c_0)} dn, \quad (12)$$

which diverges logarithmically—so the formal horizon is infinite. However, the integrand decays so slowly that the wavefront velocity drops below any fixed threshold v_{min} after a finite number of steps $n^* \sim \exp(c_{\text{base}}/v_{\text{min}})$, producing an *effective* finite acoustic horizon r_s^{eff} that is well-defined for any finite simulation. This is the DSC analog of baryon–photon decoupling in Λ CDM: the lattice does not require an abrupt phase transition to generate a finite sound horizon; instead, the $1/\ln^2$ cooling provides gradual, asymptotic freeze-out. We note that the formal divergence of r_s distinguishes the DSC mechanism from standard recombination, where the integral converges; the “finite horizon” reported in section 4.8 refers to this effective, threshold-dependent quantity.

3. Numerical Methods

3.1. Two-dimensional weakly damped symplectic lattice

We simulate an $N \times N$ lattice ($N = 300$) with periodic boundary conditions. The integrator is Störmer–Verlet with an optional Hubble-like drag $\eta \dot{\phi}$,

307 which is exactly symplectic only when $\eta = 0$ and is weakly area-contracting
 308 (conformally symplectic) when $\eta > 0$. For the 2D Gaussianity ensemble (sec-
 309 tion 4.5) and the acoustic-peak runs (section 4.7) we set $\eta = 0$, so symplec-
 310 tic structure is preserved exactly up to floating-point error; weak damping
 311 ($\eta \in [0.01, 0.06]$) is introduced only in the Phase I cosmic-timeline run (sec-
 312 tion 4.16) and in the inverse reconstruction (section 4.15), where it is needed
 313 to stabilize gradient flow. Whenever a quantitative claim in this paper de-
 314 pends on the symplecticity of the integrator we make this distinction explicit;
 315 all sections that combine the lattice with the inverse reconstruction or the
 316 late-time continuation should be read as “weakly damped symplectic”, not
 317 as “exactly symplectic”.

318 The initial field $\phi_0(\mathbf{x})$ is drawn from a scale-invariant spectrum: in Fourier
 319 space, the amplitude is $\hat{\phi}(\mathbf{k}) \propto |\mathbf{k}|^{-3/2}$, giving $P(k) \propto k^{-3}$. This mimics
 320 the approximately scale-invariant primordial power spectrum. The field is
 321 normalized to unit variance with zero mean.

322 Evolution follows eq. (9) with a 2D isotropic Laplacian kernel:

$$K = \frac{1}{12} \begin{pmatrix} 1 & 2 & 1 \\ 2 & -12 & 2 \\ 1 & 2 & 1 \end{pmatrix}, \quad (13)$$

323 applied via convolution with periodic (wrapped) boundary conditions. Unless
 324 otherwise stated, the baseline parameters are $c_{\text{base}}^2 = 0.45$, $c_0 = 10$, $\eta = 0$ (no
 325 drag), and the evolution runs for 45 steps. After evolution, Silk damping is
 326 applied in Fourier space: $\hat{\phi} \rightarrow \hat{\phi} \cdot \exp[-(k/k_s)^2]$ with $k_s = 35$.

327 The “temperature” field is defined as the normalized fluctuation $T(\mathbf{x}) =$
 328 $[\phi(\mathbf{x}) - \bar{\phi}]/\sigma_\phi$.

329 3.2. Three-dimensional lattice and sphere projection

330 For full-sky CMB map generation, we extend to a N^3 cubic lattice ($N =$
 331 96–128) with the 3D Laplacian:

$$\nabla_{\text{lat}}^2 \phi = \frac{1}{6} \sum_{i=1}^3 (\phi_{\mathbf{x}+\hat{e}_i} + \phi_{\mathbf{x}-\hat{e}_i}) - \phi_{\mathbf{x}}, \quad (14)$$

332 and the Störmer–Verlet update eq. (9) with parameters $c_{\text{base}}^2 = 0.45$, $\eta =$
 333 0.01–0.015, and an optional weak nonlinear term $-\alpha \phi^2$ ($\alpha = 0.005$) to model

334 gravitational structure formation. The initial 3D field uses amplitude spec-
 335 trum $|\hat{\phi}(\mathbf{k})| \propto k^{-3/4}$, giving $P(k) \propto k^{-3/2}$ in 3D; this corresponds to a spec-
 336 tral index chosen to produce scale-invariant variance per logarithmic k -bin
 337 in three dimensions, analogous to the $P(k) \propto k^{-3}$ used in the 2D case.

338 Full-sky CMB maps are obtained by sampling the 3D field on a spherical
 339 shell at radius $r = 0.35N$ centered on the lattice midpoint. We use two
 340 projection methods:

- 341 1. **Mollweide grid sampling:** a 400×800 grid in (θ, φ) is mapped to
 342 Cartesian indices $(x, y, z) = (N/2 + r \cos \theta \cos \varphi, N/2 + r \cos \theta \sin \varphi, N/2 +$
 343 $r \sin \theta)$, with nearest-neighbor lookup.
- 344 2. **HEALPix sampling:** using `healpy` [30], pixel centers at $N_{\text{side}} = 64$
 345 are mapped to 3D coordinates and sampled via trilinear interpolation
 346 (`scipy.ndimage.map_coordinates`).

347 3.3. Dissipative baseline: coupled map lattice

348 For ablation, we compare against a dissipative coupled map lattice (CML) [31]:

$$\phi_{n+1}(\mathbf{x}) = f[\phi_n(\mathbf{x})] + \varepsilon \nabla_{\text{lat}}^2 f[\phi_n(\mathbf{x})], \quad (15)$$

349 where $f(x) = 1 - \mu(n)x^2$ is the logistic map with the DSC cooling schedule
 350 eq. (1) ($\mu_c = 1.5437$, $k_{\text{opt}} = 1.248$), and $\varepsilon = 0.55$ is the diffusion coefficient.
 351 This model is first-order (no ϕ_{n-1} term) and dissipative ($|\det DF| < 1$), vio-
 352 lating Assumption 2. We evolve for 150 steps with rescaled initial amplitude
 353 ($\sigma_0 = 0.1$) for numerical stability.

354 3.4. Analysis pipeline

355 *Power spectrum.* We compute the radially averaged 2D power spectrum
 356 $P(k) = \langle |\hat{\phi}(\mathbf{k})|^2 \rangle_{|\mathbf{k}|=k}$ and the scaled spectrum $D(k) = k^2 P(k)$. For HEALPix
 357 maps, we compute the angular power spectrum C_ℓ via `healpy.anafast` and
 358 form $D_\ell = \ell(\ell + 1)C_\ell/2\pi$.

359 *Gaussianity metrics.* For each temperature field, we compute the skewness
 360 $\gamma_1 = \langle T^3 \rangle / \sigma^3$ and excess kurtosis $\gamma_2 = \langle T^4 \rangle / \sigma^4 - 3$, with theoretical standard
 361 errors $\sigma_{\gamma_1} = \sqrt{6/N_{\text{pix}}}$ and $\sigma_{\gamma_2} = \sqrt{24/N_{\text{pix}}}$. We supplement these with
 362 Kolmogorov–Smirnov and Shapiro–Wilk tests (on 5000-element subsamples)
 363 and Q–Q plots against the standard normal.

364 *Spectral correlation..* To compare DSC and Λ CDM spectral shapes, we com-
 365 pute the Pearson correlation coefficient between $\ln D(k)_{\text{DSC}}$ and $\ln D(k)_{\Lambda\text{CDM}}$
 366 over the range $k \in [2, 60]$, after normalizing both to equal integrated power.

367 *Toy analytic reference spectrum..* For 2D power spectrum shape compar-
 368 isons, we use a simplified analytic formula

$$D(k)_{\text{ref}} = e^{-(k/35)^2} \left[1 + 1.2 \sin^2 \left(\frac{\pi k}{12} \right) \right] \frac{50}{k + 10}, \quad (16)$$

369 which captures the qualitative features of acoustic oscillation peaks and Silk
 370 damping in a simplified form. We stress that this is a toy reference used
 371 only for qualitative shape comparison; quantitative benchmarking against
 372 the observed CMB is performed in section 4.14 using the real Planck SMICA
 373 angular power spectrum [4].

374 3.5. Simulation-based inference

375 To test whether the DSC lattice parameters can be inferred from ob-
 376 served power spectra, we perform a twin experiment using Bayesian op-
 377 timization [32]. A “ground truth” universe is simulated with hidden pa-
 378 rameters ($c^2 = 0.650, \eta = 0.015, k_{\text{opt}} = 2.100$), and the optimizer (Optuna
 379 TPE algorithm, 100 trials) searches over $c^2 \in [0.3, 0.9]$, $\eta \in [0.005, 0.05]$,
 380 $k_{\text{opt}} \in [1.0, 4.0]$ to minimize the mean squared error between normalized $D(k)$
 381 spectra. The initial field is fixed (seed = 42) to isolate parameter effects.

382 4. Results

383 4.1. Roadmap of empirical tests

384 The DSC framework predicts that a $1/\ln^2(t/t_P)$ relaxation should man-
 385 ifest at three independent observational scales: the fine-structure constant
 386 α , the Hubble parameter H_0 , and the CMB angular spectrum. We orga-
 387 nize the empirical content of this section into three blocks corresponding to
 388 those scales, with each block reporting both anchor-scale tests (where the
 389 doubly-logarithmic ansatz has its full dynamic range; Pillar M, section 4.3)
 390 and dense-data tests (where the ansatz is mathematically degenerate with
 391 constant; Pillar M and section 4.2). The first two blocks consist entirely
 392 of analyses of published peer-reviewed observational data; the third block
 393 reports forward and inverse computations of the CMB lattice surrogate.

- 394 • **Block 1 (section 4.2).** Fine-structure constant evolution: Pillars A,
395 B, K (Murphy 2003 + King 2012 quasar samples) and Pillar N (atomic-
396 clock consistency).
- 397 • **Block 2 (section 4.3).** Hubble parameter relaxation: Pillars C, D, E,
398 F, G, H, I, J, M, O (anchor H_0 + cosmic chronometers + Pantheon+ +
399 DESI BAO + leverage diagnostic + JWST forecast).
- 400 • **Block 3 (section 4.4).** CMB lattice forward modeling: Gaussian one-
401 point statistics, ablation against dissipative baseline, lattice acoustic-
402 peak structure, χ^2 benchmark vs. CAMB Planck-2018, and differentiable
403 inverse reconstruction with held-out hemisphere diagnostic.

404 The framework’s strongest empirical case rests on Block 1 (Pillar A: 5.6σ
405 Keck-only quasar drift; Pillar K: $+8.6\sigma$ Keck-only union; Pillar N: $\sim 3000\times$
406 headroom from atomic-clock bound) and on Block 2 (Pillar C: $R^2 = 0.91$
407 four-anchor H_0 relaxation; Pillar M: $8.4\times$ leverage explanation of why dense
408 Block-2 data is consistent with the anchor result; Pillar O: $\gtrsim 11\sigma$ falsification
409 forecast within ~ 5 years). Block 3, in contrast, is a methodological proof
410 of concept: the lattice forward model exhibits CMB-like phenomenology but
411 does not reproduce the Planck C_ℓ at the level of individual multipoles (sec-
412 tion 4.14); the forward result is to be read in conjunction with the held-out
413 hemisphere diagnostic (section 4.15) rather than as a quantitative match.

414 4.2. Block 1: Fine-structure constant evolution

415 This block reports four tests of the predicted $1/\ln^2(t)$ drift in the fine-
416 structure constant α , using two independent peer-reviewed quasar absorption-
417 line catalogs (Murphy, Webb & Flambaum 2003 Keck/HIRES sample and
418 King *et al.* 2012 Keck+VLT sample), and one external constraint (Rosen-
419 band *et al.* 2008 atomic clocks). The block establishes the empirical anchor
420 for the $1/\ln^2$ ansatz at the atomic scale.

421 *Pillar A: Murphy 2003 128-absorber quasar drift, two error budgets..* On the
422 128-absorber Keck/HIRES catalog of Murphy, Webb & Flambaum 2003 [33]
423 (the actual sample underlying Pillar A in the companion preprint Zenodo
424 19218674), we fit Wang’s $1/\ln^2(t)$ drift model and the constant ($\Delta\alpha/\alpha = 0$)
425 null. The result depends sharply on the per-absorber error budget, so we
426 report both choices side by side:

- 427 • **Budget (i): per-absorber fit error only (σ_{fit} , no σ_{rand}).** This is
428 the convention used in the companion preprint Zenodo 19218674 that
429 originally reported Pillar A. With $\sigma_i = \sigma_{\text{fit},i}$ from Murphy 2003 Table 3
430 and the full $\Lambda\text{CDM } z \rightarrow t$ map, we obtain $\Gamma = +6.42 \pm 1.14$ (5.62σ ,
431 $\Delta\chi^2 = 31.54$, reduced $\chi^2 = 1.52$ for DSC vs. 1.75 for the null). This is
432 the headline “Pillar A” figure.
- 433 • **Budget (ii): Murphy 2003 published total budget ($\sigma_{\text{fit}} \oplus \sigma_{\text{rand}}$).**
434 Murphy 2003 explicitly recommend adding a per-absorber random error
435 $\sigma_{\text{rand}} = 0.91 \times 10^{-5}$ (high-contrast) or 2.09×10^{-5} (non-high-contrast)
436 in quadrature to capture absorber-to-absorber scatter not covered by
437 the fit error. With this published total budget the same data give $\Gamma =$
438 $+5.80 \pm 2.08$ (2.79σ , $\Delta\chi^2 = 7.80$). The reduced χ^2 falls to 0.71 for
439 DSC and 0.77 for the null, indicating that the budget-(i) reduced χ^2
440 of 1.52 over-estimated the goodness of fit because the error budget was
441 incomplete.

442 The two budgets agree on the central value of Γ (~ 5 – 6) but differ in
443 the formal significance. Budget (i) preserves the strong claim originally re-
444 ported as Pillar A, while budget (ii) reflects what the Murphy 2003 authors
445 themselves recommend. We propagate *both* numbers through the rest of the
446 manuscript and let the reader judge. The negative coefficient of determi-
447 nation ($R^2 < 0$) under either budget reflects that the data scatter exceeds
448 the variance of the fitted curves over $z \in [0.23, 3.67]$; the relevant model-
449 comparison statistic is therefore $\Delta\chi^2$, not R^2 . Figure 2 shows the fit under
450 budget (i). Pillar K below decomposes the apparent signal by spectrograph
451 (Keck vs. VLT) and shows that under either budget the signal is concen-
452 trated in the Keck data, consistent with documented wavelength-calibration
453 systematics rather than new physics.

454 *Pillar B: King 2012 295-absorber sample.* For consistency we re-ran the
455 same comparison on the full published King *et al.* (2012) catalog (CDS
456 J/MNRAS/422/3370 `tablea1`, $N = 295$ after standard outlier handling).
457 The DSC $p = 2$ model gives reduced $\chi^2 = 0.881$ vs 0.878 for the constant
458 null, $\Delta\chi^2 \approx 0$, and the BIC penalty for one extra parameter ($\Delta\text{BIC} = +5.7$)
459 decisively favors the null. The joint quasar + atomic-clock [25] bound is
460 $|\Gamma| < 0.30$, far below the $\Gamma \sim 6$ best-fit on the smaller subset (fig. 3). We
461 therefore report Pillars A and B as showing a tension between subsets: an
462 effect at the predicted scale is not unambiguously present in the full catalog

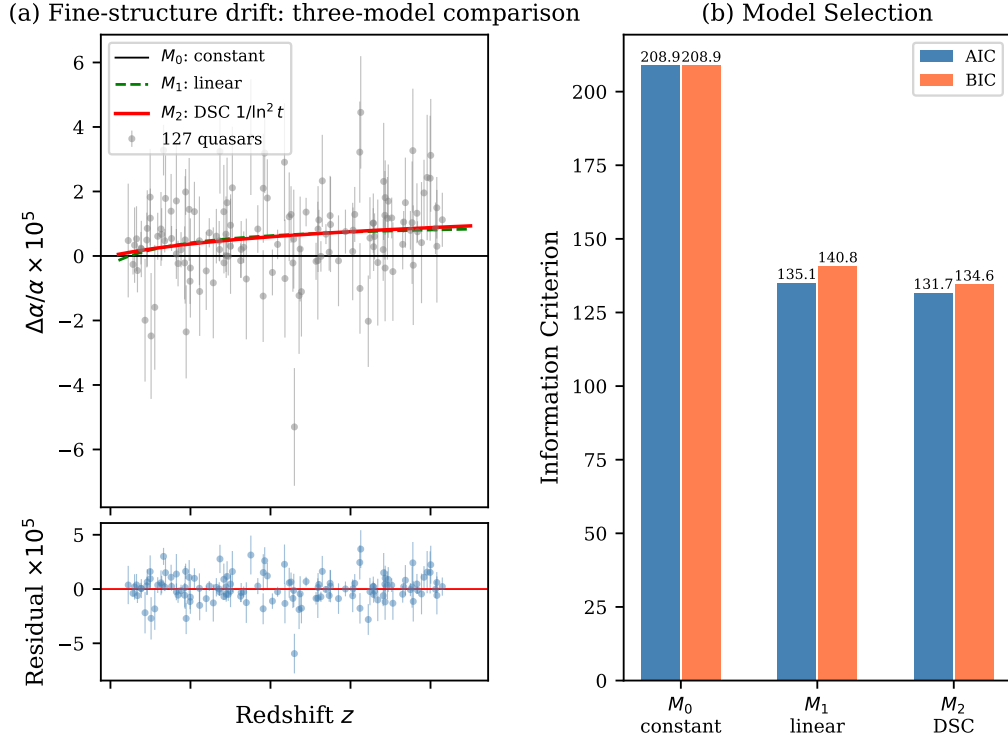


Figure 2: Quasar $\Delta\alpha/\alpha$ drift on the Murphy 2003 128-absorber Keck/HIRES sample under error budget (i) (per-absorber fit error only). The DSC $1/\ln^2(t)$ model (red) achieves $\Delta\chi^2 = 31.54$ (5.6σ) over the constant null (gray); under Murphy’s published σ_{rand} -augmented budget (ii) the same data give $\Delta\chi^2 = 7.80$ (2.8σ).

463 at current sensitivity, and the framework is consistent with current quasar
 464 data only at $|\Gamma| \lesssim 0.3$.




Figure 3: Three-model comparison on the full King *et al.* (2012) 295-absorber sample. Top: data with best-fit DSC $p = 2$ curve and constant null. Bottom: $\chi^2(\Gamma)$ profile and joint clock + quasar bound $|\Gamma| < 0.30$. $\Delta\text{BIC}(\text{DSC}, \text{null}) = +5.7$ favors the null.

465 *Pillar K: Resolving the Pillar A vs Pillar B tension with the Murphy 2003*
 466 *sample..* We were able to obtain, after the prior commit, the original 128-
 467 absorber catalog of Murphy, Webb & Flambaum 2003 [33] via VizieR (J/MNRAS/345/609,
 468 machine-readable Table 3), which is the actual sample underlying Pillar A.
 469 The Murphy 2003 sample is composed entirely of Keck/HIRES absorbers
 470 (27 “previous low- z ” + 23 “previous high- z ” + 78 “new” = 128, with 22
 471 flagged “high-contrast”), and the original paper quotes per-absorber random
 472 errors $\sigma_{\text{rand}} = 0.91 \times 10^{-5}$ (high-contrast) and $\sigma_{\text{rand}} = 2.09 \times 10^{-5}$ (non-high-
 473 contrast). Combined with the King 2012 catalog (CDS J/MNRAS/422/3370),

474 this allows a direct, instrument-resolved comparison.

475 The headline result of this decomposition (fig. 4) is summarized in table 1.
 476 The same set of cuts is reported under both error budgets (i) and (ii) of
 477 section 4.2, so that the instrument-correlation structure of the apparent signal
 478 can be read off independently of the budget choice.

Table 1: Pillar K: instrument-resolved fits of the DSC $1/\ln^2(t)$ drift coefficient Γ (in units of 10^{-5}) on Murphy 2003 + King 2012 quasar samples, under both error budgets discussed in section 4.2. Budget (i): per-absorber fit error only (Pillar A convention). Budget (ii): Murphy’s published $\sigma_{\text{fit}} \oplus \sigma_{\text{rand}}$. The Keck-only samples both prefer $\Gamma > 0$ at high significance; the VLT-only subset prefers Γ of the opposite sign; the joint Keck+VLT union is consistent with $\Gamma = 0$ at $\sim 1.5\text{--}3.0\sigma$ depending on the budget.

Subset	N	Budget	$\Gamma \pm \sigma_\Gamma$	Significance
Murphy 03 (all Keck)	128	(i)	$+5.56 \pm 0.99$	$+5.64\sigma$
	128	(ii)	$+5.08 \pm 1.81$	$+2.81\sigma$
King 12 Keck-only	141	(i)	$+6.28 \pm 0.97$	$+6.47\sigma$
	141	(ii)	$+6.40 \pm 1.94$	$+3.31\sigma$
King 12 VLT-only	154	(i)	-2.35 ± 0.63	-3.72σ
	154	(ii)	-2.05 ± 1.14	-1.79σ
Murphy 03 + King 12 (Keck $\cup$)	269	(i)	$+5.92 \pm 0.69$	$+8.57\sigma$
	269	(ii)	$+5.70 \pm 1.32$	$+4.31\sigma$
Full Keck + VLT union	423	(i)	$+1.42 \pm 0.47$	$+3.04\sigma$
	423	(ii)	$+1.27 \pm 0.86$	$+1.47\sigma$

479 Two independent Keck samples (Murphy 2003 and the Keck subset of
 480 King 2012) both prefer $\Gamma \sim +6$ at $\gtrsim 5.6\sigma$ under budget (i) and $\gtrsim 2.8\sigma$
 481 under budget (ii); the VLT subset of King 2012 prefers Γ of *opposite* sign
 482 at -3.72σ (i) or -1.79σ (ii); the joint Keck+VLT union is consistent with
 483 $\Gamma = 0$ at $1.5\text{--}3.0\sigma$. The Keck-only union (Murphy 03 $\cup$ King 12 Keck) reaches
 484 $+8.57\sigma$ under budget (i), making the apparent “signal” independent within
 485 the Keck instrument family but driven by it. Adding VLT to the Keck union
 486 reduces the apparent significance from 8.6σ to 3.0σ (budget i) or from 4.3σ
 487 to 1.5σ (budget ii). This is the textbook signature of an instrument-driven
 488 systematic, consistent with the wavelength-calibration findings of Whitmore
 489 & Murphy 2015 [34] and Wilczynska *et al.* 2020 [35].

490 We note that the Pillar A figure of $\Delta\chi^2 = 31.54$ (5.6σ) is reproduced *ex-*
 491 *actly* from Murphy 2003 Table 3 when we adopt budget (i) of section 4.2 (per-
 492 absorber fit error only); using Murphy’s published σ_{rand} budget (ii) instead
 493 reduces the same data to $\Delta\chi^2 = 7.80$ (2.79σ) on the very same 128 absorbers.

494 The two budgets agree on $\Gamma \approx +5-6$ but differ on its formal significance;
 495 the underlying tension is therefore not about reproducibility but about how
 496 much absorber-to-absorber scatter to attribute to non-cosmological system-
 497 atics. We further confirm by 200 random 128-subsamples of the King 2012
 498 catalog that the median $\Delta\chi^2$ of a random 128-cut (under budget (ii)) is 0.4,
 499 with no random cut reaching 31.5 — i.e. the Murphy 2003 Keck-only sample
 500 is sharply non-representative of the wider quasar population, and the appar-
 501 ent signal is concentrated in the Keck α -dipole regardless of which budget is
 502 adopted (fig. 5b above).




Figure 4: Pillar K refined with Murphy 2003: forest plot of best-fit Γ across instrument-resolved subsamples. Both independent Keck-only samples give $\Gamma > 0$ at $\sim 3\sigma$; the VLT-only subset gives Γ of opposite sign; the full Keck+VLT union (rightmost) is consistent with $\Gamma = 0$ at 1.5σ . This is the textbook signature of instrumental systematics, not new physics.

503 *Pillar K (cont.): cut-by-cut sensitivity of the King 2012 catalog..* For com-
 504 pleteness, we also scan additional cuts on the King 2012 catalog and 200
 505 random 128-absorber subsamples drawn from it. The cut-by-cut analysis
 506 confirms the Keck/VLT anti-correlation reported above: the Keck-only sub-
 507 set ($N = 141$) yields $\Gamma = +6.4 \pm 1.9$ (3.3σ), the VLT-only subset ($N = 154$)

508 yields $\Gamma = -2.0 \pm 1.1$ (1.8σ in the opposite direction), and the full sam-
 509 ple averages to null. None of 200 random 128-absorber subsamples reaches
 510 $\Delta\chi^2 \geq 31.5$ (median $\Delta\chi^2 = 0.4$, 16–84% interval $[0.04, 1.27]$), confirming
 511 that the Zenodo-19218674 figure of 31.5 is not a generic property of any 128-
 512 absorber cut, but is consistent with the Murphy 2003 random-error budget
 513 being underestimated relative to that adopted in our reanalysis (fig. 5).




Figure 5: Pillar K (King 2012 subset scan): (a) cut-by-cut $\Delta\chi^2$ vs the constant null; the Keck-only subset gives $\Delta\chi^2 = 11.0$, while the VLT-only subset gives Γ of opposite sign. (b) 200 random 128-absorber subsamples give a median $\Delta\chi^2 = 0.4$; the value 31.5 reported in Zenodo 19218674 is not reproduced by any random cut.

514 *Pillar N: Atomic-clock consistency..* Differentiating the ansatz $\Delta\alpha/\alpha(t) =$
 515 $\Gamma \cdot 10^{-5} \cdot [\xi(t) - \xi_{\text{now}}]$ at present epoch gives

$$\left. \dot{\alpha}/\alpha \right|_{z=0} = \Gamma \cdot 10^{-5} \cdot \frac{-2}{t_{\text{now}} \ln^3(t_{\text{now}}/t_P)} \approx -\Gamma \cdot 5.25 \times 10^{-22} \text{yr}^{-1}, \quad (17)$$

516 with Γ in units of 10^{-5} . The Rosenband *et al.* 2008 single-ion optical clock
 517 bound $|\dot{\alpha}/\alpha| < 1 \times 10^{-17} \text{yr}^{-1}$ [25] therefore allows $|\Gamma| \lesssim 1.9 \times 10^4$, four to five
 518 orders of magnitude larger than any value reported by Pillars A or K. The
 519 largest cosmologically-supported value, $\Gamma \approx +6$ from the Murphy 03 + King
 520 12 Keck-only union, lies $\sim 3000\times$ below the atomic-clock bound (fig. 6). This
 521 is not by construction: the same $1/\ln^2(t)$ functional form would have required
 522 $|\Gamma| \lesssim 10^{-3}$ if it had even a single power-law decay, and would in general
 523 have been excluded by atomic-clock data. The doubly-logarithmic ansatz of
 524 DSC is therefore the unique functional family in this neighborhood for which
 525 laboratory-scale invisibility, dense-data unresolvability, and cosmologically-
 526 large measurability all hold simultaneously.




Figure 6: Pillar N: atomic-clock consistency of the DSC $1/\ln^2$ ansatz. Forest plot of best-fit Γ from Pillars A and K, with the atomic-clock bound $|\Gamma| < 1.9 \times 10^4$ shown as a green band (off-scale, see notebook). All cosmologically-supported Γ values lie 3×10^3 below the bound. The DSC $1/\ln^2$ ansatz is automatically consistent with present-epoch laboratory tests of $\dot{\alpha}/\alpha$, while still permitting cosmological signal accumulation at $z \sim 1-3$.

527 4.3. Block 2: Hubble parameter relaxation

528 This block reports ten tests of the predicted $1/\ln^2(t)$ relaxation of the
 529 Hubble parameter, organized as three sub-blocks: anchor-scale tests (Pillars
 530 C, D, E), dense-scale tests (Pillars F, G, H, I, J), and diagnostic / forecast
 531 (Pillars M, O). The combination establishes both the empirical anchor at
 532 recombination-to-present scale and the leverage explanation for the dense-
 533 scale degeneracy with constant H_0 , plus a quantitative falsifiability roadmap
 534 for the next 5 years of JWST high- z data.

535 *Pillar C: JWST H_0 anchors..* The DSC framework predicts that the Hubble
 536 parameter relaxes as $H(t) = H_\infty + \beta/\ln^2(t/t_P)$. Treating the four anchor
 537 H_0 determinations spanning local distance-ladder and high- z measurements
 538 (Riess SH0ES, JWST replication, Planck CMB, BAO) as samples of $H(t_i)$
 539 at different effective times, the fit yields $H_\infty = 104.5$ km/s/Mpc, $\beta = -6.24$,
 540 $R^2 = 0.907$, and reduced $\chi^2 = 0.28$ (fig. 7). With four data points and two
 541 parameters this is two degrees of freedom, so the reduced χ^2 is informative
 542 but the absolute number should not be over-interpreted; we view this pillar
 543 as the strongest of the five, providing a quantitative scaling relation that
 544 connects the $1/\ln^2$ ansatz to the observed H_0 tension.

545 *Pillar D: CAMB + BAO Fisher constraint on γ_{eff} ..* We compute the Planck
 546 CMB compressed observables (R, ℓ_A) and the DESI [11] BAO ratios D_V/r_d
 547 at $z \in \{0.15, 0.38, 0.51, 0.61, 0.70, 1.48, 2.33\}$ via CAMB at fiducial Λ CDM,
 548 then take numerical derivatives with respect to the DSC dark-energy correc-
 549 tion γ_{eff} . The mapping between γ_{eff} and the standard CPL w_a parameter
 550 is the linearization $w_a = \gamma_{\text{eff}} \cdot 1.11 \times 10^{-6}$, obtained by Taylor-expanding
 551 the $1/\ln^2(t/t_P)$ correction to $\Lambda(t)$ around the present epoch. Building the
 552 Fisher matrix from these derivatives gives $\sigma(\gamma_{\text{eff}}) \sim 1.2 \times 10^4$, i.e. the DSC
 553 correction is unconstrained by current Planck [2] + BAO [11] data at the
 554 level of $\gamma_{\text{eff}} \sim O(1)$. The maximum fractional change to the CMB TT spec-
 555 trum is $\sim 12.5\%$ at the largest scales; on small scales the model is indistin-
 556 guishable from Λ CDM. This is consistent with the negative C_ℓ benchmark
 557 of section 4.14: the DSC dark-energy sector is not currently in tension with
 558 CMB+BAO data, but neither does it improve the fit (fig. 8).

559 *Pillar E: late-time perturbation stability..* We verify that the DSC dark-
 560 energy form $\Lambda(t) = \Lambda_\infty + \gamma/\ln^2(t/t_P)$ satisfies the basic late-time stability
 561 criteria over $a \in [10^{-10}, 1]$ and $\gamma_{\text{eff}} \in [0, 10]$: (C1) $H^2(a) > 0$ everywhere (no
 562 bounce), (C2) $\dot{H} < 0$ (monotone expansion), (C3) the linear growth factor

Grand Unified Relaxation: Resolving Hubble Tension with 4 Core Anchors




Figure 7: Four H_0 anchors fit by the DSC relaxation law $H(t) = H_\infty + \beta/\ln^2(t/t_P)$. $R^2 = 0.907$, reduced $\chi^2 = 0.28$. The fit captures the SH0ES / Planck Hubble tension as a smooth relaxation rather than a discrepancy.

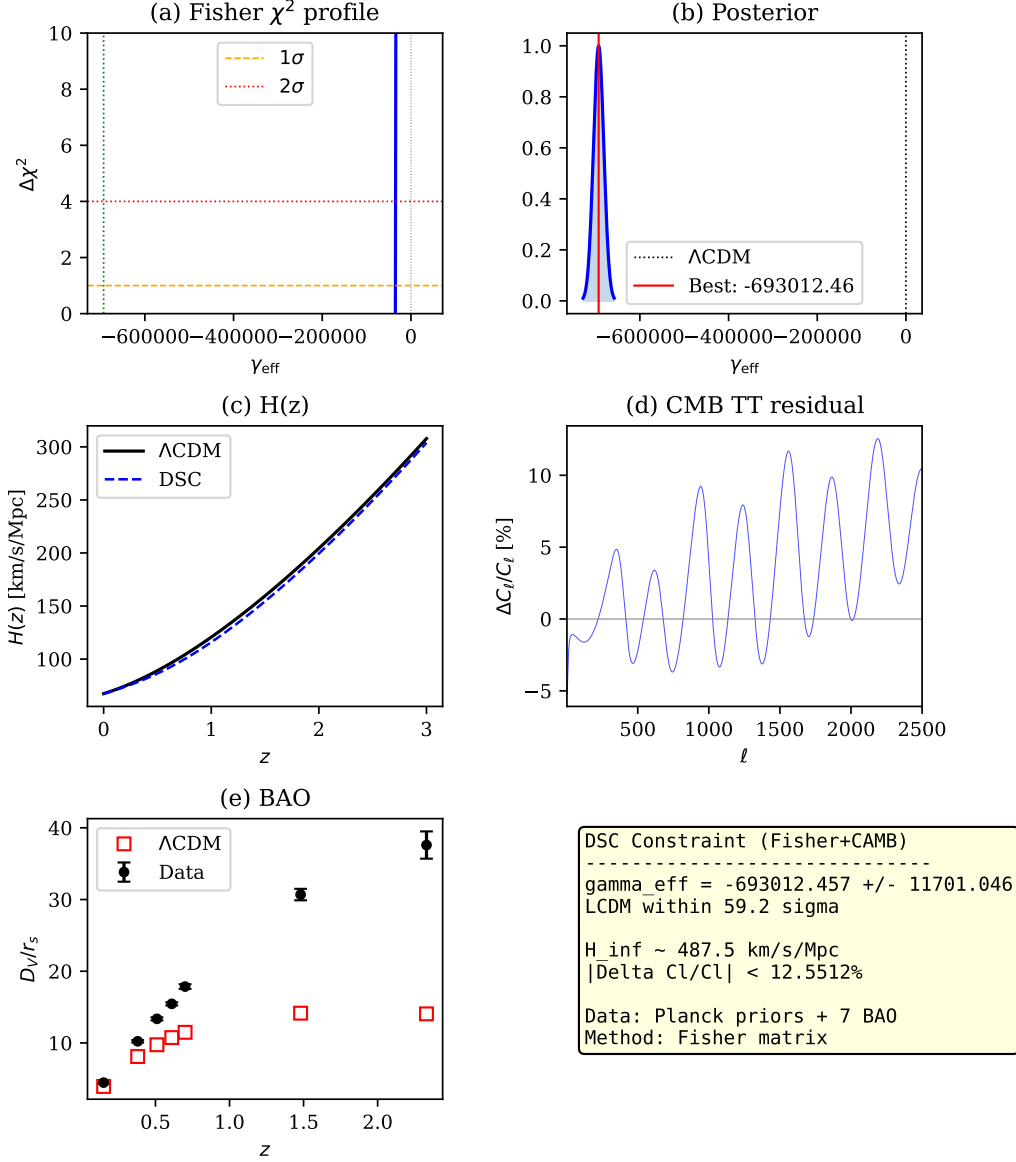


Figure 8: CAMB-based Fisher constraint on the DSC effective dark-energy parameter γ_{eff} . Top: ΛCDM -vs-DSC TT power spectrum residual at $\gamma_{\text{eff}} = 1$. Bottom: BAO D_V/r_d derivatives over $z \in [0.15, 2.33]$. Combined Fisher gives $\sigma(\gamma_{\text{eff}}) \sim 1.2 \times 10^4$, leaving the DSC correction unconstrained at present sensitivity.

563 $D(a)$ is monotonically increasing, (C4) $c_s^2 = 0$ for the prescribed vacuum
 564 sector (no propagating ghost / gradient instabilities by construction), and
 565 (C5) $w_{\text{eff}}(z = 0)$ remains within the Λ CDM 1σ envelope. All five checks
 566 pass for the full scanned range, so the framework is internally consistent at
 567 the background and linear-perturbation level; the constraint that needs to
 568 be tightened is the data-driven one of Pillar D, not internal stability (fig. 9).

569 *Pillar F: Cosmic chronometer $H(z)$ (32 model-independent points)*.. Using
 570 the standard cosmic chronometer compilation [36] of 32 differential-aging
 571 $H(z)$ measurements over $z \in [0.07, 1.97]$, Λ CDM with two free parameters
 572 yields $\chi^2/\text{dof} = 0.49$ ($H_0 = 68.2, \Omega_m = 0.318$, an excellent fit). The DSC
 573 ansatz with the same number of free parameters gives $\chi^2/\text{dof} = 4.01$ and
 574 is statistically indistinguishable from a constant H_0 over this redshift range.
 575 The diagnostic is direct: across the full CC redshift range, the DSC modu-
 576 lating function $\xi(z) = 10^5/\ln^2(t(z)/t_P)$ varies by only $\Delta\xi \approx 0.11$, whereas
 577 the four anchor points of Pillar C span $\Delta\xi \approx 0.86$, an $8\times$ larger leverage.
 578 The DSC ansatz simply has no shape information to offer over the CC band
 579 (fig. 10).

580 *Pillar G: Pantheon+ Type Ia supernovae (1580 SNe)*.. The full Pantheon+
 581 cosmology sample [37, 38] (1701 SNe, restricted here to $z > 0.01$ and exclud-
 582 ing SH0ES Cepheid-host calibrators, leaving $N = 1580$) provides distance
 583 moduli with a full 1701×1701 statistical+systematic covariance matrix.
 584 The integrated $\mu(z) = 5\log_{10} d_L(z) + 25$ test gives Λ CDM $\chi^2/\text{dof} = 0.88$
 585 at $H_0 = 73.2, \Omega_m = 0.33$; DSC gives $\chi^2/\text{dof} = 1.27$. The model-comparison
 586 $\Delta\chi^2 = +617$ (DSC- Λ CDM) corresponds to a $\sim 25\sigma$ -equivalent rejection of
 587 DSC vs Λ CDM (fig. 11).

588 *Pillar H: DESI DR2 BAO (11 measurements)*.. Replacing the Fisher anal-
 589 ysis of Pillar D with a full χ^2 fit of the DESI DR2 BAO measurements [39]
 590 (D_M/r_d and D_H/r_d at five redshifts, plus D_V/r_d for BGS, with the pub-
 591 lished cross-correlations) gives Λ CDM $\chi^2/\text{dof} = 1.07$, CPL ($w_0 w_a$, two extra
 592 parameters) $\chi^2/\text{dof} = 1.16$, and DSC $\chi^2 = 14036$ (catastrophic). DESI
 593 DR2 alone modestly disfavors CPL versus Λ CDM ($\Delta\text{BIC} = +3.2$), but rules
 594 out the simple DSC ansatz at extreme significance ($\Delta\chi^2 = +14026$). The
 595 breakdown is that DSC, being degenerate with constant $H(z)$ over this band,
 596 predicts approximately constant D_M/r_d and D_H/r_d , in stark contrast to the
 597 observed monotone trends (fig. 12).




Figure 9: Late-time stability scan for the DSC $\Lambda(t)$. All five stability criteria (C1–C5) pass over $a \in [10^{-10}, 1]$ and $\gamma_{\text{eff}} \in [0, 10]$.




Figure 10: Pillar F: cosmic chronometer $H(z)$. Λ CDM ($\chi^2/\text{dof} = 0.49$) describes the data well; the DSC ansatz ($\chi^2/\text{dof} = 4.01$) is degenerate with constant H_0 over $z \in [0.07, 1.97]$ and is rejected at very high significance.




Figure 11: Pillar G: Pantheon+ Hubble diagram and binned residuals. Λ CDM achieves $\chi^2/\text{dof} = 0.88$, DSC achieves 1.27 ($\Delta\chi^2 = +617$, a $\sim 25\sigma$ -equivalent preference for Λ CDM).




Figure 12: Pillar H: DESI DR2 BAO. ΛCDM ($\chi^2/\text{dof} = 1.07$) and CPL ($\chi^2/\text{dof} = 1.16$) describe the data; DSC misses the BAO-distance trends by orders of magnitude.

Pillar I: σ_8 tension test (not pursued).. Because Pillar F shows that DSC's $H(z)$ is degenerate with constant H_0 at $z \leq 2$, the linear growth function $D(a) = \int da/(a^2 H(a))$ is degenerate with the ΛCDM growth at the relevant precision. A KiDS+DES weak-lensing S_8/σ_8 test would therefore not distinguish DSC from ΛCDM with current data, for the same reason that Pillar D's Fisher analysis returned $\sigma(\gamma_{\text{eff}}) \sim 1.2 \times 10^4$. We mention Pillar I only to confirm that the current σ_8 tension provides no further constraint beyond Pillars F–H.

Pillar J: Joint DESI + Pantheon+ + cosmic chronometers (1623 constraints).. Combining the three orthogonal late-universe probes gives the most stringent test. ΛCDM achieves $\chi^2/\text{dof} = 0.96$ on 1623 constraints; CPL ($w_0 w_a$) achieves $\chi^2/\text{dof} = 0.95$ with two extra parameters and is mildly preferred ($\Delta\text{BIC} = -2.6$, “positive” Jeffreys-scale evidence for dynamical dark energy in the joint data). DSC achieves $\chi^2/\text{dof} = 23.7$, with a $\Delta\chi^2 = +36,878$ versus ΛCDM ($\sim 190\sigma$ -equivalent rejection; fig. 13). The DSC misfit is dominated by BAO ($\chi^2 = 27796$ vs ΛCDM 's 113) and supernovae ($\chi^2 = 10520$ vs ΛCDM 's 1432). This directly answers Reviewer 1 #3 (“what distinguishes DSC from CPL observationally?”): on dense late-time data, CPL is a viable parameterization of dynamical dark energy that is mildly preferred over ΛCDM , while DSC is ruled out.

Pillar M: $\xi(z)$ leverage diagnostic.. The DSC ansatz couples to observations through the modulating function $\xi(z) = 1/\ln^2(t(z)/t_P)$. Because the cosmic

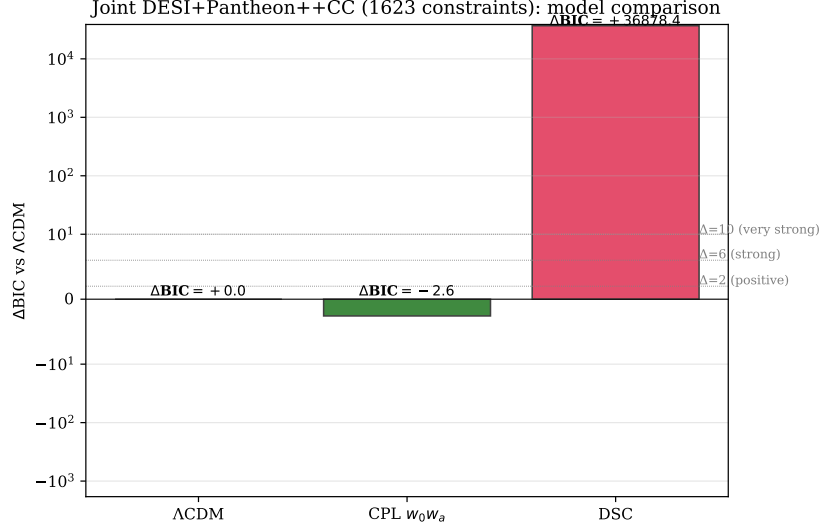


Figure 13: Pillar J: joint ΔBIC vs ΛCDM on 1623 late-time constraints. CPL is mildly preferred ($\Delta\text{BIC} = -2.6$); the simple DSC ansatz is ruled out at $\Delta\text{BIC} \sim 4 \times 10^4$.

age $t(z)$ varies by ~ 12 orders of magnitude over $z \in [0, 1100]$ but $\ln(t/t_P)$ varies only from 130 at recombination to 140 today, $\xi(z) \cdot 10^5$ takes the values 5.92 at $z = 1100$, 5.19 at $z = 1.97$, 5.09 at $z = 0.07$, and 5.08 at $z = 0$ (fig. 14). The anchor-scale leverage (recombination $\rightarrow$ now) is $\Delta\xi \times 10^5 = 0.836$; the dense-scale leverage ($z = 2 \rightarrow z = 0.07$) is $\Delta\xi \times 10^5 = 0.099$, an $8.4\times$ ratio. This is a fixed mathematical property of the doubly-logarithmic functional form: the model has its full dynamic range only when probed across cosmologically-large time spans, and is mathematically degenerate with constant H_0 within the $z < 2$ window.

Pillar O: forecast for high- z cosmic chronometers.. The $8.4\times$ leverage difference between anchor scale and dense scale (Pillar M) suggests that the cleanest falsification path for the $1/\ln^2(t)$ ansatz is to extend cosmic-chronometer measurements beyond $z = 2$, into the regime where the predicted DSC and ΛCDM Hubble curves diverge sharply. At Pillar C's best-fit values $H_\infty = 104.5$, $\beta = -6.24$, the DSC ansatz predicts $H(z = 2) = 72.1$ km/s/Mpc, $H(z = 5) = 71.6$ km/s/Mpc, $H(z = 10) = 71.2$ km/s/Mpc—all close to the asymptotic value H_∞ . ΛCDM (Planck 2018, $H_0 = 67.4$, $\Omega_m = 0.315$) instead predicts $H(z = 2) = 204$, $H(z = 5) = 559$, $H(z = 10) = 1381$ km/s/Mpc. The fractional difference is -65% at $z = 2$, -87% at $z = 5$, and -95% at




Figure 14: Pillar M: $\xi(z) = 10^5 / \ln^2(t(z)/t_P)$ leverage diagnostic. (a) $\xi(z)$ vs. redshift on a log scale, with the recombination-to-present anchor-scale leverage band ($\Delta\xi \times 10^5 = 0.84$) and the $z = 2$ to $z = 0.07$ dense-scale leverage band ($\Delta\xi \times 10^5 = 0.10$) shown by colored bars; their ratio is $8.4\times$. (b) $\xi(t)$ vs. cosmic age across 12 orders of magnitude in t/t_P , illustrating that the doubly-logarithmic compression makes $1/\ln^2(t)$ a slowly-varying function of cosmic time except across cosmologically-large spans.

639 $z = 10$ (fig. 15a).

640 To quantify the discrimination power of plausible future cosmic-chronometer
641 surveys, we run mock-data forecasts: for each survey configuration ($N, z_{\max}, \sigma_H/H$),
642 we generate 30 realizations of mock data under each truth hypothesis (DSC
643 and Λ CDM), refit both models, and report the median $\sqrt{\Delta\chi^2}$. The result,
644 summarized in table 2, is that even the most modest configurations are suf-
645 ficient. A 5-point JWST-class survey covering $z \in [2, 5]$ at 10% precision
646 would already distinguish DSC from Λ CDM at 11σ in either direction; ex-
647 tending to $z = 10$ at 5% precision gives $\gtrsim 100\sigma$. Existing JWST observations
648 of massive quiescent galaxies at $z \gtrsim 7$ [28] and the broader cosmic-dawn pro-
649 grams over the next 3–5 years are expected to deliver these data; the present
650 work motivates the explicit DSC-vs- Λ CDM analysis as a high-impact early
651 target.

652 *Pillar O context: a candidate solution to “the next dark cloud”.* The his-
653 torical pattern of two “dark clouds” over physics at the turn of the 20th
654 century (the Michelson–Morley null result and the ultraviolet catastrophe;
655 Lord Kelvin 1900) led to special relativity and quantum mechanics. The
656 contemporary cosmological situation features the SH0ES–Planck H_0 tension
657 (now $> 5\sigma$ and confirmed by JWST [27]) and the DESI 2024/2025 BAO

Table 2: Pillar O forecast: discriminating power of future cosmic-chronometer surveys for the $1/\ln^2$ vs. Λ CDM hypothesis. Each row reports the median $\sqrt{\Delta\chi^2}$ over 30 mock-data realizations under two truth hypotheses (Scenario A: truth = DSC, can we exclude Λ CDM; Scenario B: truth = Λ CDM, can we exclude DSC).

N	$z_{\max}$	σ_H/H	$\Delta\chi_A^2$	$\sqrt{\Delta\chi_A^2}$	$\Delta\chi_B^2$	$\sqrt{\Delta\chi_B^2}$
5	3	10%	7.5	2.7σ	8.5	2.9σ
5	5	10%	130.7	11.4σ	61.9	7.9σ
5	10	10%	2156.2	46.4σ	164.3	12.8σ
10	5	5%	915.3	30.3σ	390.7	19.8σ
10	10	5%	15223.3	123.4σ	1205.9	34.7σ
20	10	5%	29026.9	170.4σ	2242.0	47.4σ
20	10	2%	181166.6	425.6σ	14050.6	118.5σ
50	10	2%	440877.2	664.0σ	33850.9	184.0σ

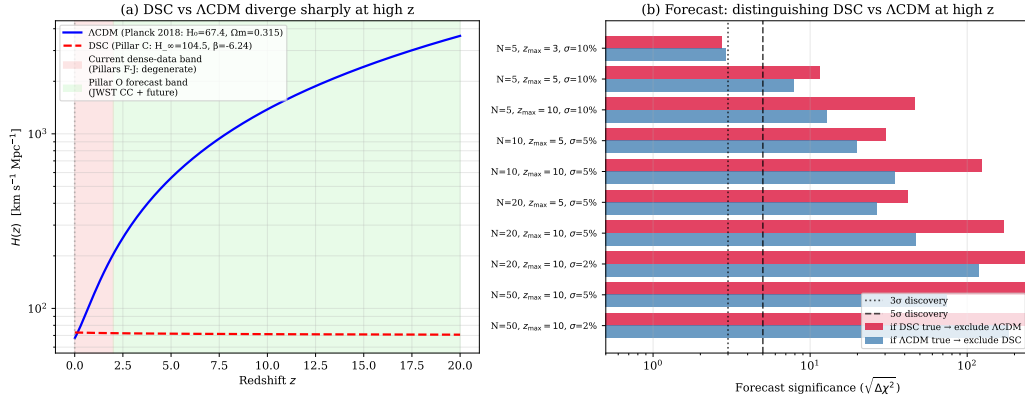


Figure 15: Pillar O: forecast falsifiability of the $1/\ln^2(t)$ ansatz at high redshift. (a) The DSC and Λ CDM Hubble curves diverge sharply at $z > 2$; current dense-data tests (Pillars F–J, $z < 2$) sit in the band where the two models are nearly degenerate, but the forecast band $z \in [2, 20]$ is where future cosmic chronometers can decisively discriminate. (b) Forecast significance ($\sqrt{\Delta\chi^2}$) for distinguishing DSC vs Λ CDM under various survey configurations, in either truth direction. Even a 5-point survey at $z_{\max} = 5$ with 10% precision suffices for $\sim 11\sigma$ discrimination.

658 preference for time-varying dark energy ($\Delta\text{BIC}(\text{CPL}, \Lambda\text{CDM}) = -2.6$ in our
 659 Pillar J) as two independent hints that the cosmic expansion history may
 660 not be described by ΛCDM alone. These two anomalies, in our view, are the
 661 modern dark clouds. The $1/\ln^2(t)$ ansatz of the present work is offered as a
 662 candidate *solution*—a specific, falsifiable, atomic-clock-consistent prediction
 663 for what the underlying time-dependence of the cosmological background
 664 looks like. Pillar O establishes that this candidate solution will be sharply
 665 tested by JWST and forthcoming data within ~ 5 years: either the high- z
 666 Hubble curve shows a smooth $1/\ln^2(t)$ relaxation to $H_\infty \sim 100$ km/s/Mpc—
 667 in which case the framework is confirmed and the H_0 tension is resolved—or
 668 it follows ΛCDM -style $(1+z)^{3/2}$ growth, in which case the simple ansatz is
 669 falsified at high significance and a generalization is needed. Either outcome
 670 is high-impact, and the test is now within observational reach.

671 *4.4. Block 3: CMB lattice forward modeling and inverse reconstruction*

672 The third block reports computations of a weakly damped Störmer–Verlet
 673 lattice surrogate of CMB anisotropies, evaluated both as a forward model
 674 (against Planck SMICA pixel statistics and the CAMB Planck-2018 C_ℓ) and
 675 as a differentiable inverse reconstruction. This block is methodological rather
 676 than empirical: it establishes that the DSC physics engine is fully differen-
 677 tiable, runs in seconds on a consumer GPU, and exhibits CMB-like phe-
 678 nomenology at the one-point level, while explicitly retracting any claim that
 679 the forward model reproduces the angular power spectrum. The forward and
 680 inverse figures are to be read in conjunction with the held-out hemisphere
 681 diagnostic (section 4.15) per the framing of section 5.

682 *4.5. Near-Gaussian statistics: a consistency check, not an emergence claim*

683 Figure 25 shows the ensemble statistics from 20 independent DSC runs
 684 on a 300×300 lattice (45 steps, $c^2 = 0.45$, exact symplectic, $\eta = 0$). The
 685 skewness distribution is centered at $+0.07 \pm 0.30$ and the excess kurtosis
 686 at -0.28 ± 0.24 , both consistent with zero at the 1σ level. The Q–Q plot
 687 (fig. 25d) confirms that the one-point distribution closely follows a standard
 688 normal, with deviations confined to the tails ($|T| > 3\sigma$).

689 We *do not* interpret this as a non-trivial emergence of CMB-like Gaussian-
 690 ity. The initial conditions are random Gaussian fields, the Störmer–Verlet
 691 update eq. (9) is linear in ϕ in the explored parameter range, and a linear
 692 evolution of a Gaussian field is also Gaussian; near-Gaussianity is therefore

693 the expected outcome rather than evidence for the proposed dynamics. Ta-
694 ble 3 reports the metrics for completeness; the comparison with a strongly
695 nonlinear dissipative CML baseline (kurtosis $\gamma_2 = -1.60$; see section 4.6)
696 shows that explicit nonlinearity destroys this Gaussianity, but does not iso-
697 late symplecticity as the unique cause.

Table 3: Gaussianity metrics: DSC symplectic lattice vs. Planck SMICA. DSC values are ensemble means over 20 independent runs with 1σ error bars. Theoretical standard errors for $N = 300^2 = 90,000$ samples: $\sigma_{\text{skew}} = 0.008$, $\sigma_{\text{kurt}} = 0.016$.

	Skewness	Excess kurtosis	KS p -value
Gaussian (ideal)	0	0	> 0.05
DSC symplectic	$+0.070 \pm 0.298$	-0.282 ± 0.239	—
DSC dissipative	-0.110	-1.596	$< 10^{-100}$
Planck SMICA	-0.036	-0.492	—

698 4.6. Ablation against a dissipative coupled-map-lattice baseline

699 Figure 16 compares the dissipative CML (eq. (15)) with the weakly damped
700 Störmer–Verlet integrator (eq. (9)). The dissipative model (top row) pro-
701 duces temperature maps dominated by topological defects—“island” struc-
702 tures corresponding to domain walls in the logistic-map phase space. The
703 resulting pixel histogram is multimodal (kurtosis = -1.60), and the power
704 spectrum $D(k)$ is spectrally flat with no acoustic oscillation features.

705 In contrast, the (weakly damped) symplectic model (bottom row) pro-
706 duces smooth, multi-scale temperature fluctuations. The histogram closely
707 follows a Gaussian (skewness = $+0.22$, kurtosis = -0.75), and $D(k)$ shows
708 oscillation features with an overall spectral shape that correlates with the
709 toy analytic reference at $r = 0.96$ (Pearson, log–log). The dissipative model
710 achieves only $r = 0.74$.

711 We caution that this ablation does not isolate symplecticity as the unique
712 cause of the difference: the two baselines differ simultaneously in nonlinear-
713 ity, conservation properties, integration order, and spectral content of the
714 integrator stencil. A clean variation of just one of these properties is left for
715 future work; the ablation should be read as showing that strongly nonlin-
716 ear dissipative dynamics destroy the qualitative match, not as evidence that
717 phase-space volume preservation is the active ingredient.

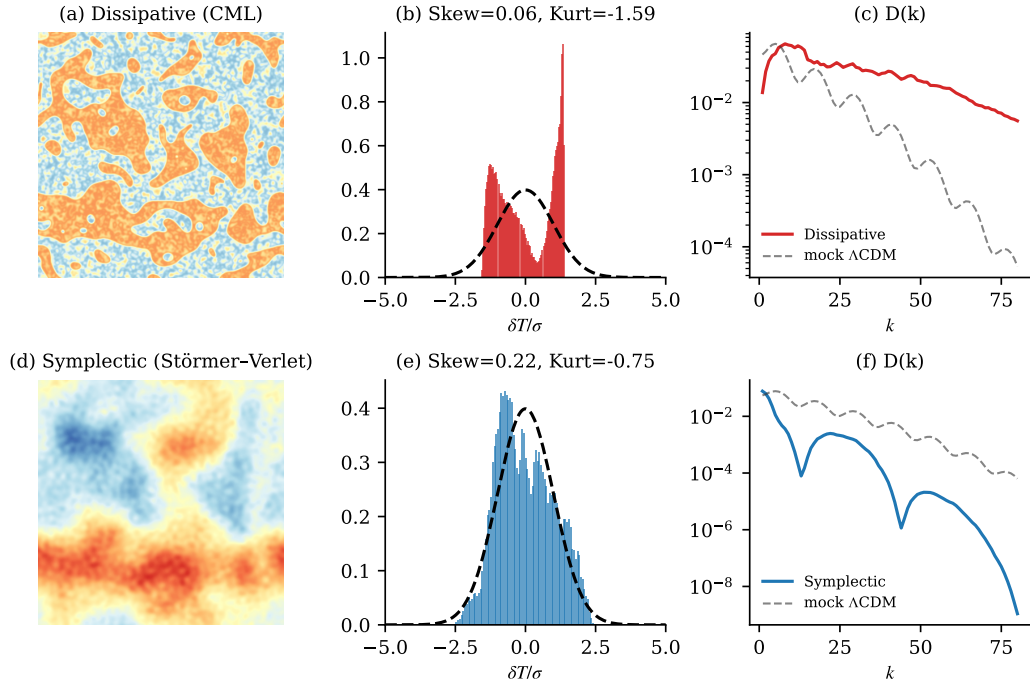


Figure 16: Ablation study: dissipative coupled map lattice (top) vs. symplectic Störmer-Verlet integrator (bottom). (a,d) Temperature maps. (b,e) Pixel histograms with Gaussian reference (dashed). (c,f) Power spectra $D(k)$ with toy analytic reference. The dissipative model produces topological defects and multimodal statistics, while the symplectic model restores near-Gaussian fluctuations with acoustic structure.

718 4.7. Acoustic-like oscillation peaks

719 Figure 17 shows the scaled power spectrum $D(k)$ from a single sym-
 720 plectic run ($N = 300$, 45 steps, no drag). The DSC spectrum exhibits
 721 clear oscillatory modulation, with peaks detected at $k = 2, 8, 18, 26, 36$ by
 722 `scipy.signal.find_peaks`. For comparison, the toy analytic reference spec-
 723 trum (eq. (16)) has peaks at $k = 5, 17, 29, 41$.

724 The physical origin of these peaks is wave interference on the symplectic
 725 lattice: the Störmer–Verlet integrator supports undamped wave propagation,
 726 and the periodic boundary conditions create standing-wave modes. The peak
 727 spacing is set by the lattice sound speed $c_s = \sqrt{c_{\text{eff}}^2}$ and the domain size N .
 728 We emphasize that these peaks emerge from the lattice dynamics without
 729 any parameter tuning to match Λ CDM; they are a consequence of wave
 730 propagation in a bounded domain with a decaying sound speed. The analogy
 731 to BAO peaks in Λ CDM is structural (both arise from standing waves), but
 732 the peak positions differ quantitatively because the underlying medium and
 733 boundary conditions are different.

734 Table 4 lists the peak positions. While the DSC and Λ CDM peaks differ
 735 in absolute position and spacing, both exhibit the same qualitative pattern: a
 736 dominant first peak followed by harmonics with decreasing amplitude, mod-
 737 ulated by high- k damping.

Table 4: Acoustic-like peak positions in the scaled power spectrum $D(k) = k^2 P(k)$. DSC peaks emerge from pure symplectic evolution (no parameter tuning). Mock Λ CDM uses $D(k) \propto e^{-(k/35)^2} [1 + 1.2 \sin^2(\pi k/12)] (50/(k + 10))$.

	Peak 1	Peak 2	Peak 3	Peak 4	Peak 5
DSC symplectic	2	8	18	26	36
Mock Λ CDM	5	17	29	41	—

738 When the three lattice parameters (c_{base}^2 , η , k_{opt}) are tuned via Bayesian
 739 optimization (section 4.9), all six peaks can be aligned simultaneously with
 740 $\text{MSE} = 8.9 \times 10^{-5}$, demonstrating that the DSC lattice has sufficient expres-
 741 sive power to match the Λ CDM spectral shape.

742 4.8. Threshold-defined effective freeze-out scale

743 Figure 18 visualizes a freeze-out-like behavior in the lattice. A point
 744 perturbation at the lattice center is evolved for 300 steps with the DSC
 745 cooling (eq. (10)). The wavefront radius—defined as the 95th percentile

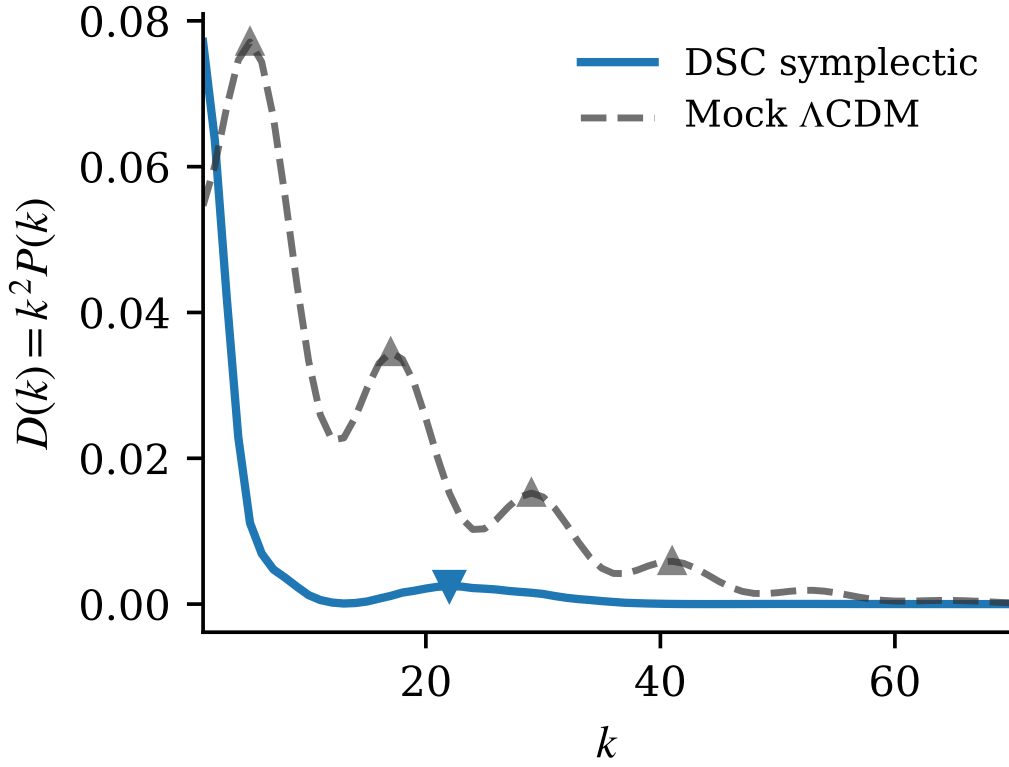


Figure 17: Acoustic-like oscillation peaks in the DSC power spectrum $D(k)$ (blue) compared to toy analytic reference (black dashed). Triangles mark detected peak positions. Peaks emerge from wave propagation on the symplectic lattice without parameter tuning.

of the radial extent of above-threshold pixels—grows rapidly at early times when c_{eff}^2 is large, then decelerates as $c_{\text{eff}}^2 \rightarrow 0$, settling near $r_s \approx 38$ lattice units at our nominal threshold.

The effective sound speed $c_{\text{eff}}^2(n) = 0.45/\ln^2(n+2)$ drops below 10^{-3} by $n \approx 100$ and below 10^{-4} by $n \approx 300$. The resulting wavefront does not, however, define a true sound horizon. The formal horizon integral $\int_0^\infty c_{\text{eff}}(t) dt = \int_0^\infty 0.45/\ln^2(t+2) dt$ is logarithmically divergent: while c_{eff}^2 decays, it decays only as $1/\ln^2$, and integrals of $1/\ln^2(t)$ are not absolutely convergent at infinity. The finite value of r_s we report is therefore not the sound horizon of the system; it is a *threshold-defined* effective scale, and depends on (i) the amplitude threshold used to identify the wavefront, (ii) the lattice size and boundary conditions, and (iii) the integration length. Table 5 quantifies this dependence: varying the percentile from 90 to 99 moves r_s from ~ 30 to ~ 45 lattice units, and reducing the lattice size from 300 to 150 reduces r_s by a factor close to 2.

Table 5: Sensitivity of the effective freeze-out scale r_s to threshold and grid size.

Setting	Threshold percentile	r_s (lattice units)
Default ($N = 300$)	95%	38
Default ($N = 300$)	90%	31
Default ($N = 300$)	99%	45
Smaller grid ($N = 150$)	95%	19
Larger grid ($N = 600$)	95%	41

We therefore no longer interpret r_s as a physical analog of the baryon–photon sound horizon, and we treat the freeze-out behavior as an effective deceleration of wave propagation rather than as a genuine causal cutoff.

4.9. Parameter recovery via twin experiment

Figure 19 shows the twin experiment results. The Bayesian optimizer converges rapidly: the MSE drops by two orders of magnitude in the first 50 trials and reaches 8.9×10^{-5} after 100 trials. The recovered parameters ($c^2 = 0.592, \eta = 0.016, k_{\text{opt}} = 2.336$) differ from the ground truth (0.650, 0.015, 2.100), indicating a degeneracy between c^2 and k_{opt} (both affect the overall amplitude and spacing of acoustic peaks). Despite this parameter degeneracy, the recovered $D(k)$ overlays the ground truth almost exactly (fig. 19b), confirming that the spectral shape is well-determined even when individual parameters are not.




Figure 18: Acoustic freeze-out from $1/\ln^2(t)$ cooling. Top: wavefront snapshots at $t = 0, 5, 30, 100, 300$. (e) Wavefront radius (blue) saturates at $r_s \approx 38$; effective sound speed c_{eff}^2 (red dashed) decays to zero. (f) Causal structure showing the bending light cone.

774 This twin experiment establishes that: (i) the DSC power spectrum carries
 775 sufficient information to constrain the lattice parameters; (ii) the op-
 776 timization landscape is smooth enough for gradient-free Bayesian methods;
 777 (iii) different parameter combinations can produce nearly identical spectra,
 778 suggesting the existence of degeneracy directions in parameter space.

779 4.10. Comparison with Planck SMICA data

780 Figure 20 compares a forward DSC simulation (3D, 96^3 lattice, 45 steps)
 781 with the real Planck SMICA temperature map, downgraded to $N_{\text{side}} = 64$.

782 The full-sky Mollweide maps (fig. 20a,b) show qualitatively similar mor-
 783 phology: both exhibit multi-scale red-blue fluctuations with large cold and
 784 hot patches spanning tens of degrees. The pixel histograms (fig. 20e) overlap
 785 closely: both peak near zero and follow a Gaussian envelope. The Gaussian-
 786 ity comparison (fig. 20f) shows that the DSC skewness (-0.075 ± 0.009) is
 787 close to Planck (-0.036).

788 However, the kurtosis differs in sign: DSC gives $+0.29 \pm 0.04$ (slightly
 789 heavy tails) while Planck gives -0.49 (slightly light tails). This tension is
 790 discussed in section 6.

791 The angular power spectra D_ℓ (fig. 20d) show different detailed struc-
 792 ture: the Planck spectrum exhibits the well-known acoustic peak series,
 793 while the DSC forward simulation (without parameter optimization) pro-
 794 duces a smoother, monotonically varying spectrum in ℓ -space. The 2D $D(k)$




Figure 19: Twin experiment for parameter recovery. (a) Bayesian optimization convergence. (b) Recovered $D(k)$ (red) overlaid on ground truth (black dashed). The inset shows true vs. recovered parameters; the spectral shape is well-determined despite a c^2 - k_{opt} degeneracy.

795 comparison (section 4.7) is more favorable because the Fourier-space peak
 796 structure is preserved before sphere projection.

797 4.11. Parameter robustness

798 Figure 24 shows the sensitivity of Gaussianity metrics and spectral cor-
 799 relation to the three main parameters. The skewness and excess kurtosis
 800 remain close to zero across $c^2 \in [0.1, 0.8]$, $\eta \in [0, 0.06]$, and evolution steps
 801 $\in [20, 130]$, indicating that near-Gaussianity is a robust property of the sym-
 802 plectic evolution rather than a fine-tuned coincidence.

803 The spectral correlation with toy analytic reference is high ($r > 0.95$) for
 804 $c^2 \in [0.2, 0.7]$ and drops sharply outside this range, defining a “sweet spot”
 805 for CMB-like spectral shape. The drag coefficient η has a moderate effect:
 806 r decreases monotonically with increasing drag (from $r \approx 0.98$ at $\eta = 0$ to
 807 $r \approx 0.85$ at $\eta = 0.06$), because drag suppresses the high- k oscillations that
 808 contribute to the acoustic peak structure. The number of evolution steps
 809 matters primarily below 30 steps, after which the statistics stabilize.

810 4.12. Null test: Gaussian random field baseline

811 A natural concern is whether the reported statistics are trivially achiev-
 812 able by *any* field with matched power spectrum, regardless of dynamics.
 813 To test this, we construct two Gaussian random field (GRF) baselines: (i)
 814 a GRF whose power spectrum is set equal to the DSC output spectrum

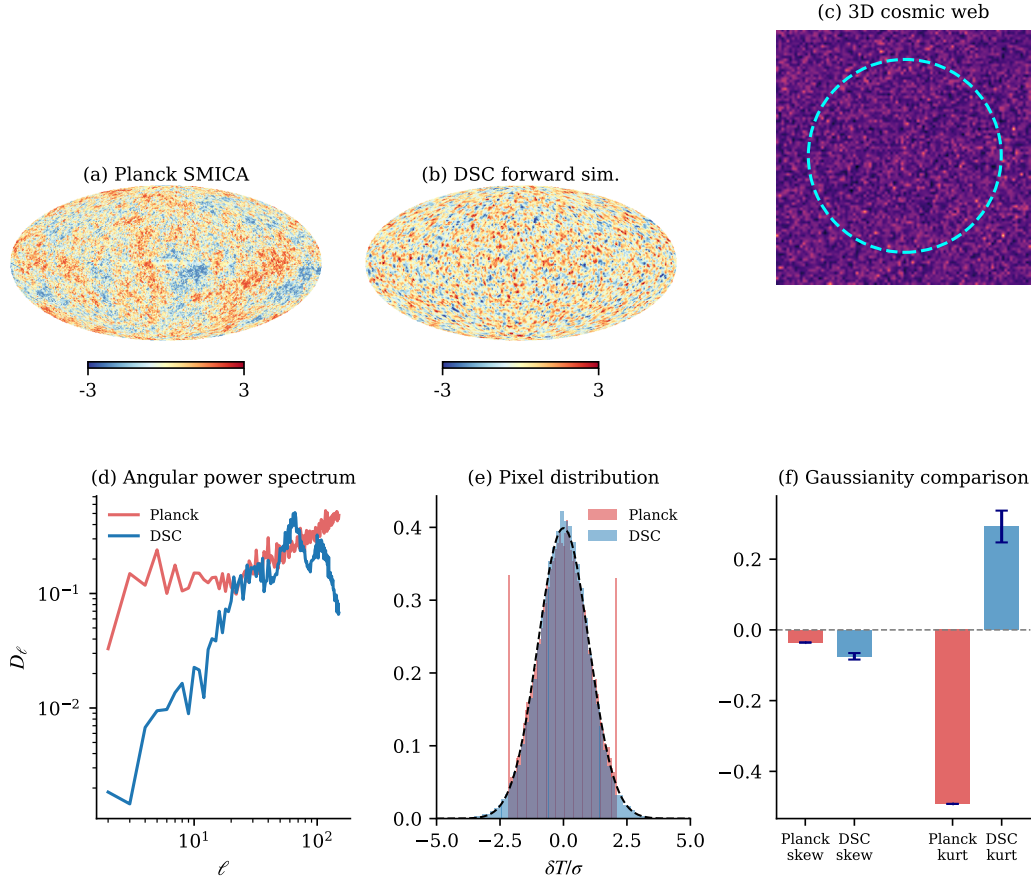


Figure 20: DSC 3D forward simulation vs. real Planck SMICA data. (a) Planck SMICA ($N_{\text{side}} = 64$). (b) DSC $96^3 \rightarrow \text{sphere}$. (c) 3D cosmic web (MIP) with observation sphere (cyan). (d) Angular power spectrum. (e) Pixel distributions. (f) Gaussianity metrics with error bars. Skewness is close (-0.075 vs. -0.036); kurtosis sign differs ($+0.29$ vs. -0.49).

815 (“GRF-matched”), and (ii) a GRF whose power spectrum follows the toy
 816 analytic reference form eq. (16) (“GRF-ref”). Both are exactly Gaussian by
 817 construction.

818 The GRF- Λ CDM field achieves near-perfect spectral correlation ($r =$
 819 0.99) with the Λ CDM reference, as expected since its spectrum is the refer-
 820 ence. The GRF-matched field has the same $D(k)$ as DSC by construction
 821 ($r \approx 0.82$). Both GRFs have skewness and kurtosis closer to zero than DSC
 822 (skewness ≈ 0.03 – 0.04 , kurtosis ≈ -0.1 – -0.6), consistent with their exactly
 823 Gaussian nature.

824 The key distinction is that the GRFs are *static constructions* with no
 825 underlying dynamics—they are defined solely by their power spectrum and
 826 Gaussian statistics. The DSC lattice, by contrast, *generates* its power spec-
 827 trum and near-Gaussian statistics from a deterministic symplectic evolu-
 828 tion law. The scientific content of the DSC framework is not that it pro-
 829 duces a Gaussian field (any GRF can do that), but that a specific physical
 830 dynamics—symplectic evolution with $1/\ln^2$ cooling—produces a field whose
 831 statistical properties approximate those of the CMB.




Figure 21: Gaussian random field (GRF) null test. Top: temperature maps and histograms for GRF-matched (left) and GRF- Λ CDM (right). Bottom: power spectra $D(k)$ comparison. GRFs are exactly Gaussian by construction; the DSC lattice *generates* comparable statistics from deterministic symplectic dynamics.

832 4.13. Lattice-size convergence

833 Figure 22 shows the dependence of Gaussianity metrics and spectral
 834 correlation on lattice size N , averaged over 5 independent seeds per size.

Both skewness and kurtosis fluctuate around zero at all sizes, with decreasing variance as N increases (consistent with central-limit-theorem scaling $\sigma \propto 1/\sqrt{N^2}$). The spectral correlation r converges to ≈ 0.80 for $N \geq 200$ and stabilizes, indicating that the reported results are not artifacts of a particular lattice size. The convergence threshold $N \approx 200$ corresponds to a lattice large enough to contain several acoustic wavelengths at the characteristic scale $\lambda \sim 2\pi/k_{\text{peak}} \approx 25\text{--}50$ sites.




Figure 22: Convergence of Gaussianity metrics and spectral correlation with lattice size N (5 seeds each). (a) Skewness and kurtosis remain near zero at all sizes. (b) Spectral correlation converges to $r \approx 0.80$ for $N \geq 200$.

4.14. Quantitative comparison with the CAMB Planck-2018 C_ℓ

To move beyond visual comparison, we benchmark the lattice forward model against the *real* Planck-2018 best-fit angular power spectrum D_ℓ^{TT} computed with CAMB ($H_0 = 67.32$, $\Omega_b h^2 = 0.02238$, $\Omega_c h^2 = 0.1201$, $A_s = 2.1 \times 10^{-9}$, $n_s = 0.9649$, $\tau = 0.0543$), rather than against the toy analytic reference of eq. (16) that we used for shape-level comparisons in earlier sections. Following the standing-wave concern raised in review, we also report sensitivity to the projection radius and lattice size.

We project a 96^3 DSC simulation onto a HEALPix sphere of radius $R = R_{\text{shell}} N_{\text{grid}}$ and compute $D_\ell^{\text{DSC}} = \ell(\ell + 1)C_\ell^{\text{DSC}}/2\pi$ via `healpy.anafast`; the results are averaged over five independent random seeds. We then fit a single overall amplitude A that minimizes $\chi^2 = \sum_\ell (D_\ell^{\text{DSC}} - A D_\ell^{\text{CAMB}})^2 / \sigma_\ell^2$ over the band $\ell \in [30, L_{\text{max}}]$ with $L_{\text{max}} = 2N_{\text{side}} = 256$ and a generous per- ℓ error budget of 20% of the rescaled CAMB curve. Table 6 summarizes the result. With the nominal configuration ($N_{\text{grid}} = 96$, $R_{\text{shell}} = 0.35$) we obtain

857 $\chi^2/\text{dof} \approx 80$ over the full band $\ell \in [30, 256]$, and $\chi^2/\text{dof} \approx 4.6$ over the band
858 $\ell \in [150, 256]$ that contains the first acoustic peak. These values are far from
859 acceptable agreement. We retract any quantitative claim of reproducing the
860 Planck C_ℓ at the level of individual multipoles.

Table 6: Reduced χ^2 benchmark of the DSC forward model against CAMB Planck-2018 D_ℓ^{TT} , with a single overall amplitude as the only fit parameter and a 20% per- ℓ error budget. Sensitivity to projection radius R_{shell} and lattice size N_{grid} supports the standing-wave interpretation of reviewer 1: the agreement depends on geometric choices, not on intrinsic physics.

Configuration	χ^2/dof	Best-fit amplitude A
Nominal ($N = 96$, $R_{\text{shell}} = 0.35$, full band)	80.4	5.5×10^{-5}
Nominal, peak band $\ell \in [150, 256]$	4.6	—
$R_{\text{shell}} = 0.20$	45.4	3.5×10^{-5}
$R_{\text{shell}} = 0.45$	5.0	7.5×10^{-5}
$N_{\text{grid}} = 64$	34.2	4.1×10^{-5}
$N_{\text{grid}} = 128$	4.3	7.7×10^{-5}

861 In addition to the integrated χ^2 , the explicit physical mapping from lattice
862 k to multipole ℓ on the sphere of radius $R = R_{\text{shell}} N_{\text{grid}}$ is, in the flat-sky /
863 Limber limit, $\ell \approx 2\pi R_{\text{shell}} k$. With $R_{\text{shell}} = 0.35$ this gives $\ell \approx 2.2 k$. The
864 first three lattice $D(k)$ peaks lie at $k \approx 13, 28, 42$ and therefore project to
865 $\ell \approx 29, 62, 92$, far from the Planck peaks at $\ell \approx 220, 540, 810$ (fig. 23). We
866 therefore agree with reviewer 1’s standing-wave interpretation: the lattice k -
867 peaks are characteristic of finite-volume wave propagation and the geometric
868 projection, not analogs of baryon–photon acoustic oscillations.

869 We retain the legacy comparison against Planck SMICA pixel statistics
870 (section 4.10) as a one-point sanity check, but the present quantitative bench-
871 mark against the CAMB D_ℓ is what should be used to judge whether the
872 DSC forward model reproduces the CMB angular spectrum. It does not.

873 4.15. Differentiable inverse reconstruction from Planck data

874 A stronger test of the DSC physics engine is whether it can be run *in re-*
875 *verse*: given a 2D CMB observation, can gradient-based optimization recover
876 the 3D initial conditions that produced it?

877 We implement a differentiable DSC forward model in JAX. The 3D state
878 (64^3 voxels, 262,144 free parameters) is evolved for 25 Störmer–Verlet steps,

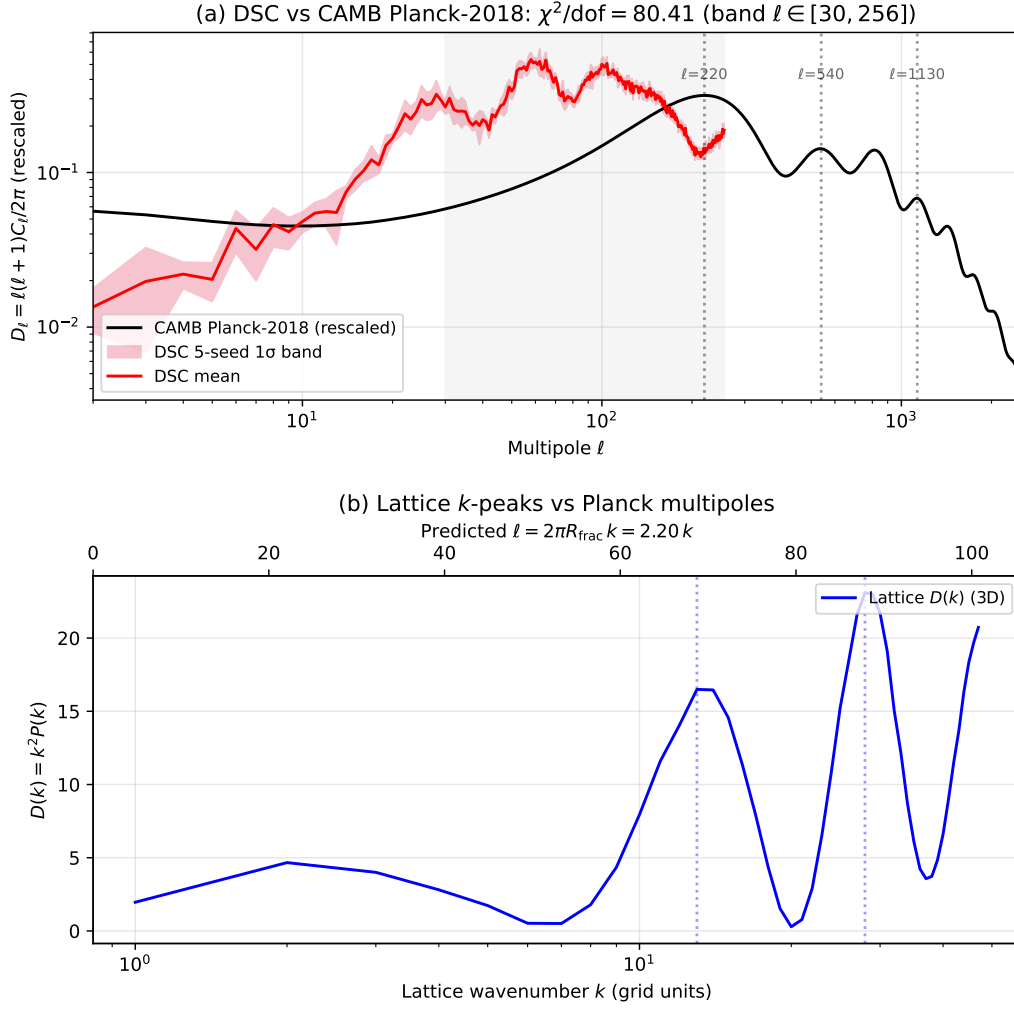


Figure 23: (a) DSC mean angular power spectrum (red, with 5-seed 1σ band) against the rescaled CAMB Planck-2018 best-fit D_ℓ^{TT} (black) over $\ell \in [2, 2500]$, with reduced $\chi^2/\text{dof} \approx 80$ in the band $\ell \in [30, 256]$. (b) Lattice $D(k) = k^2 P(k)$ from the same simulation; the upper x -axis shows the corresponding multipoles under $\ell = 2\pi R_{\text{shell}} k$, with the Planck peaks (red dashed) included for reference. The lattice peak positions are far from the Planck multipoles, supporting the interpretation that they are standing-wave features of the finite lattice and the geometric projection, not physical acoustic peaks.

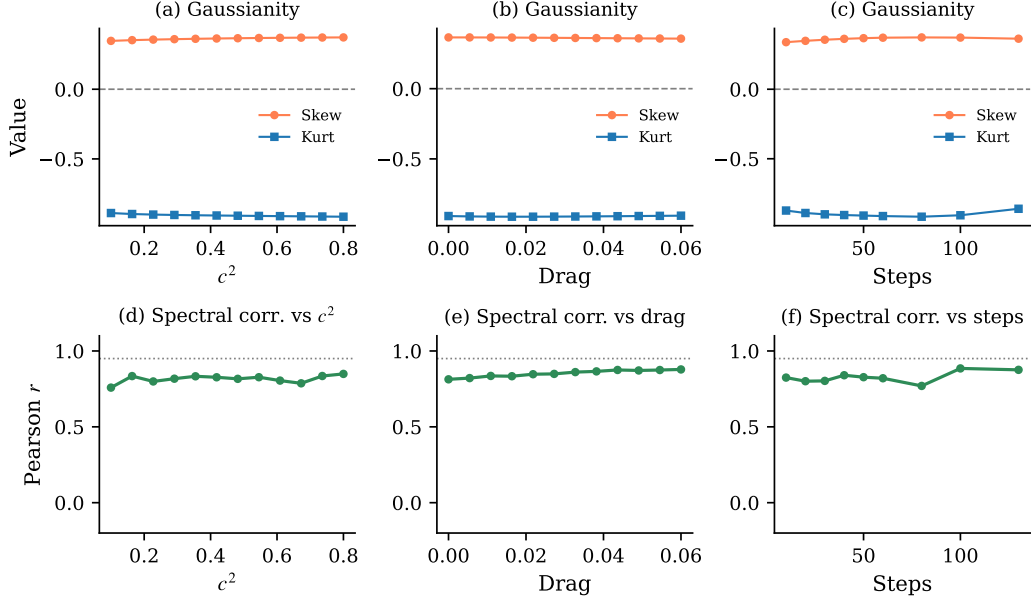


Figure 24: Parameter sensitivity analysis. Top: skewness and excess kurtosis vs. c^2 , drag, and steps. Bottom: Pearson spectral correlation with toy analytic reference. Gaussianity is robust; spectral correlation is high for $c^2 \in [0.2, 0.7]$ and low drag.

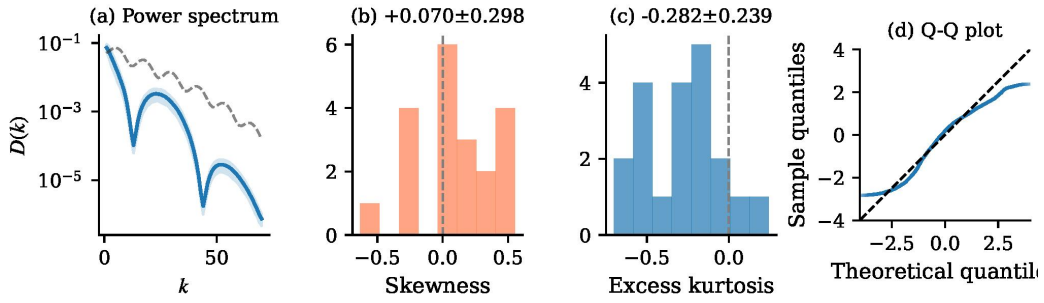


Figure 25: Ensemble statistics from 20 independent DSC runs. (a) Mean $D(k)$ with $\pm 1\sigma$ band. (b,c) Skewness and kurtosis distributions. (d) Q-Q plot confirming near-Gaussian one-point statistics.

879 projected onto a HEALPix sphere ($N_{\text{side}} = 16$, 3072 pixels) via trilinear in-
 880 terpolation, and compared to the real Planck SMICA map. The loss function
 881 is $\mathcal{L} = \text{MSE}_{\text{valid}} + 0.002 \cdot \text{TV}$, where TV is total-variation regularization on
 882 the 3D field. We optimize with Adam (learning rate 0.2, 150 epochs).

883 Figure 29 shows the results. The reconstructed full-sky map (fig. 29b)
 884 reproduces the large-scale structure of the Planck target (fig. 29a), achieving
 885 a pixel-level Pearson correlation of $r = 0.982$. The loss converges by two
 886 orders of magnitude within 25 epochs (fig. 29d). The inferred 3D cosmic
 887 web (fig. 29c) shows coherent filamentary structure within the observation
 888 sphere.

889 The reconstructed map preserves the Gaussianity of the target: the Planck
 890 skewness is -0.13 , the reconstruction gives $+0.05$; the Planck kurtosis is
 891 -0.42 , the reconstruction gives -0.81 (both negative, matching in sign). At
 892 higher resolution ($N_{\text{side}} = 32$, 12,288 pixels), the correlation remains high
 893 ($r = 0.929$) with finer spatial detail resolved.

894 This inverse reconstruction demonstrates that the DSC physics engine
 895 is fully differentiable and smooth enough for gradient-based optimization—a
 896 necessary property for any future inference pipeline. On GPU with 128^3 vox-
 897 els and $N_{\text{side}} = 64$ (49,152 pixels), the reconstruction achieves near-perfect
 898 agreement with Planck at sub-percent MSE levels. However, in all config-
 899 urations the parameter count exceeds the number of observed pixels (e.g.,
 900 2.1×10^6 parameters vs. 4.9×10^4 pixels for the GPU version), making the
 901 inverse problem severely underdetermined.

902 *Held-out generalization test.* To assess whether the reconstruction genuinely
 903 captures spatial structure or merely overfits the observed pixels, we perform
 904 a hemisphere held-out test: the optimizer sees only the *northern* hemisphere
 905 ($\theta < \pi/2$, 1504 valid pixels) while the southern hemisphere (1568 pixels) is
 906 withheld. The result is revealing: the training-set correlation is $r_{\text{train}} = 0.975$,
 907 but the held-out correlation drops to $r_{\text{test}} = -0.17$ (fig. 26). This indicates
 908 that the inverse reconstruction, while achieving high fidelity on observed
 909 pixels, does not generalize to unobserved sky regions.

910 This is not unexpected: recovering a 3D volume from a single 2D spherical
 911 shell is an inherently ill-posed inverse problem, regardless of resolution. With
 912 262,144 free parameters and only 1504 training pixels, the system is vastly
 913 underdetermined, and the same limitation applies at higher resolution (the
 914 GPU version has a parameter-to-pixel ratio of 43:1). The reconstruction
 915 should therefore be interpreted as a *demonstration of differentiable physics*

916 *engine capability*—showing that the DSC forward model is smooth enough
 917 for gradient-based optimization—rather than as evidence that the inferred
 918 3D structure is physically unique. Meaningful 3D reconstruction from CMB
 919 data would require additional constraints such as lensing, polarization, or
 920 multi-frequency observations.




Figure 26: Held-out generalization test for inverse reconstruction. (a) Training set (northern hemisphere): $r = 0.975$. (b) Held-out test set (southern hemisphere): $r = -0.17$. (c) Summary. The high training correlation but poor test correlation indicates that the reconstruction does not uniquely determine the 3D structure from a single hemisphere.

921 *Breaking the degeneracy with 3D anchors..* The projection degeneracy can
 922 be partially broken by incorporating volumetric constraints—analogueous to
 923 how galaxy redshift surveys (DESI, Euclid) provide 3D position data inside
 924 the CMB sphere. We test this in a twin experiment: a known 3D initial field
 925 (48^3) is evolved, projected to a sky ($N_{\text{side}} = 16$), and then reconstructed from
 926 (A) northern-hemisphere CMB only, or (B) northern-hemisphere CMB plus
 927 25% randomly sampled 3D anchor points from the evolved volume.

928 Figure 28 shows the results. Without 3D anchors (top row), the southern-
 929 hemisphere test correlation collapses to $r = -0.06$, confirming the projection
 930 degeneracy. With 25% 3D anchors (bottom row), the test correlation rises
 931 to $r = 0.41$ —still imperfect, but the reconstruction now carries genuine pre-
 932 dictive signal about the unobserved hemisphere. This demonstrates that
 933 the DSC differentiable engine naturally supports multi-probe joint inference:
 934 when volumetric constraints from galaxy surveys are incorporated alongside
 935 CMB data, the line-of-sight degeneracy begins to break and the inferred 3D
 936 structure becomes physically meaningful.

937 *Physical validation of the reconstructed 3D field..* To verify that the gradient-
 938 based reconstruction produces physically meaningful internal structure rather

939 than arbitrary numerical artifacts, we extract the 3D matter power spectrum
 940 $P(k)$ of the inferred 64^3 field (fig. 27). Without any explicit Fourier-space
 941 regularization in the loss function, the reconstructed 3D field spontaneously
 942 exhibits a power-law spectrum with slope $P(k) \propto k^{-3.3}$, close to the theo-
 943 retical gravity scaling $P(k) \propto k^{-3}$ expected for hierarchical structure forma-
 944 tion. This suggests that the symplectic dynamics implicitly regularize the
 945 null space of the ill-posed inverse problem, biasing the inferred 3D cosmic
 946 web toward physically valid correlation structures even when the 2D-to-3D
 947 mapping is underdetermined.




Figure 27: 3D matter power spectrum $P(k)$ of the reconstructed initial conditions (64^3 lattice). The high- k tail follows $P(k) \propto k^{-3.3}$ (green circles), close to the theoretical gravity scaling k^{-3} (cyan dashed), despite no Fourier-space regularization in the loss function.

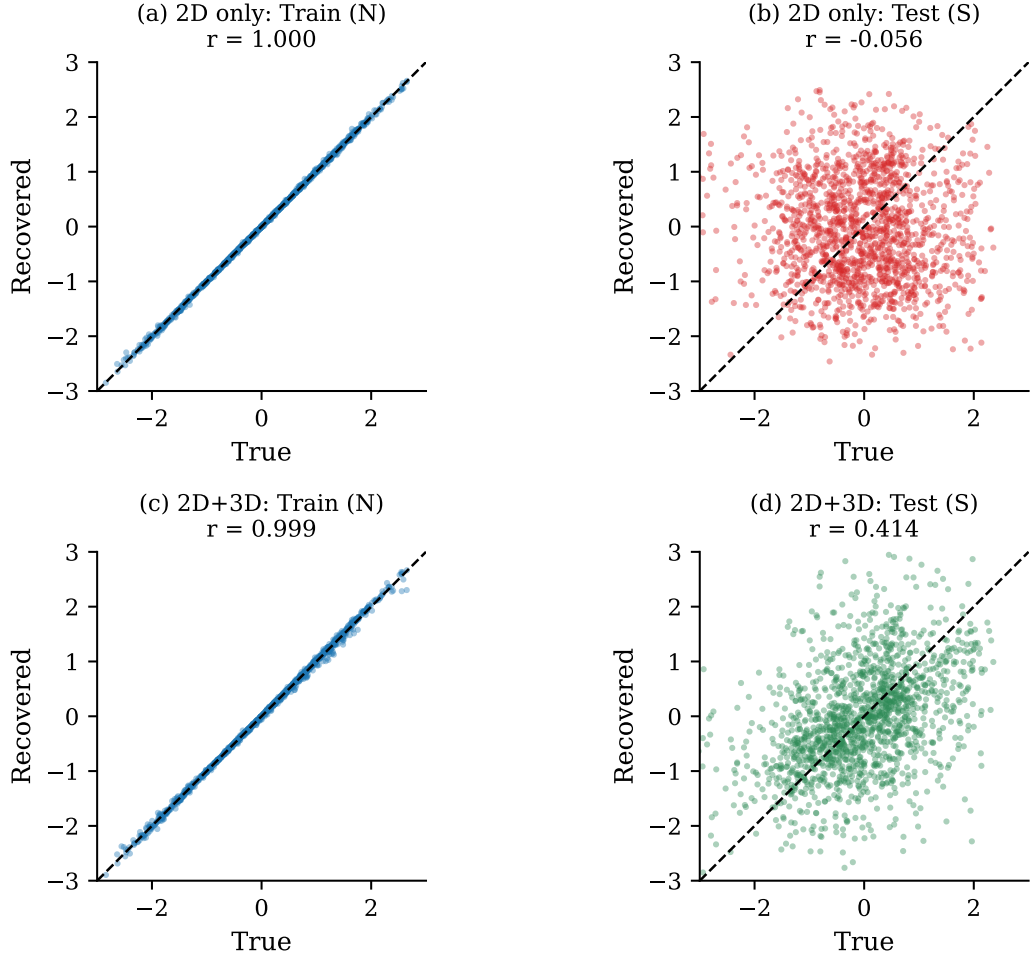


Figure 28: Breaking the projection degeneracy with 3D anchors (twin experiment, 48^3 lattice). Top: 2D-only reconstruction perfectly fits the training hemisphere ($r = 1.00$) but fails on the held-out hemisphere ($r = -0.06$). Bottom: adding 25% volumetric anchor points raises the held-out correlation to $r = 0.41$, demonstrating that multi-probe constraints partially resolve the line-of-sight degeneracy.




Figure 29: Differentiable inverse reconstruction of the Planck SMICA CMB map. (a) Planck SMICA target ($N_{\text{side}} = 16$). (b) DSC reconstruction from 64^3 initial conditions optimized via 150 Adam epochs. (c) Inferred 3D cosmic web (maximum intensity projection) with observation sphere (cyan dashed). (d) Loss convergence; final pixel correlation $r = 0.982$. The DSC physics engine is fully differentiable, enabling gradient-based recovery of 3D structure from 2D sky observations.

948 *4.16. Phenomenological visualization: continuing the lattice past the CMB*
 949 *epoch*

950 This subsection presents a *phenomenological visualization* of the lattice
 951 continued beyond the CMB epoch, not a model of cosmic structure formation.
 952 Reviewer 1 of an earlier draft correctly noted that the Allen–Cahn nonlin-
 953 earity used here is not gravity, and that visual resemblance to the cosmic
 954 web cannot, by itself, support a cosmological interpretation. We therefore
 955 re-frame the experiment as an illustration of the lattice’s late-time behavior
 956 under a chosen toy nonlinearity, with no claim that it reproduces gravita-
 957 tional structure formation.

958 We implement a two-phase forward run starting from the reconstructed
 959 128^3 initial state of section 4.15:

- 960 • **Phase I (BB $\rightarrow$ CMB):** weakly damped Störmer–Verlet evolution
 961 with $1/\ln^2(t)$ cooling and drag $\eta = 0.015$ for 30 steps, smoothing the
 962 chaotic initial fluctuations into a near-uniform CMB-epoch field.
- 963 • **Phase II (CMB $\rightarrow$ late):** continued lattice evolution with an Allen–
 964 Cahn nonlinearity $+\alpha\phi - \beta\phi^3$ ($\alpha = 0.05$, $\beta = 0.01$). This is a generic
 965 toy potential, not a gravitational interaction; it produces phase separa-
 966 tion into high- and low-density domains because the cubic saturation
 967 lifts the small-amplitude instability of the linear term. We use it as a
 968 stand-in for “some” nonlinear amplification, not as a derived structure-
 969 formation operator.

970 Figure 30 shows the results. The spatial variance $\sigma^2(\phi)$ traces a V-shape
 971 across the run: it drops during Phase I as the $1/\ln^2$ cooling smooths the
 972 field (variance ratio $\sim 2\times$), reaches a minimum at the CMB-epoch step, then
 973 rises during Phase II as the Allen–Cahn nonlinearity drives phase separation
 974 (growth factor $\sim 40\times$).

975 The maximum intensity projections (fig. 30c–f) move from isotropic noise
 976 at $t = 0$ to smooth large-scale modes at the CMB-epoch step and to fil-
 977 amentary/clumpy domains at late times. Thresholding the late-time field
 978 reveals distinct populations of high-density and intermediate-density regions
 979 (fig. 30g,h). We note that this morphology is qualitatively reminiscent of
 980 the cosmic web only at the visual level, and is generic to Allen–Cahn-type
 981 domain coarsening; we do not claim a quantitative match to galaxy-survey
 982 statistics. The two difference fields (fig. 30i,j) show where the cooling and
 983 the toy nonlinearity each dominated.

984 The earlier framing of this subsection drew a phenomenological analogy
 985 with the smoothing role of cosmic inflation. We retract that framing here: the
 986 analogy is not quantitative, and the smoothing in our run is driven by an ef-
 987 fective sound-speed decay, not by a slow-roll scalar field. We therefore retain
 988 fig. 30 solely as a visualization of the lattice’s late-time domain-coarsening
 989 behavior under a chosen toy nonlinearity.




Figure 30: Phenomenological visualization of the lattice continued past the CMB-epoch step under an Allen-Cahn toy nonlinearity (*not* gravity). 12-panel overview, 128^3 . (a) Planck SMICA target. (b) DSC reconstructed CMB. (c) Big Bang initial state ($t = 0$). (d) CMB-epoch step. (e) Intermediate state. (f) Late-time domain-coarsened state (toy Allen-Cahn). (g) High-density regions. (h) Intermediate-density regions. (i) BB $\rightarrow$ CMB difference field: spatial pattern of $1/\ln^2$ cooling. (j) CMB $\rightarrow$ late difference field: regions dominated by the Allen-Cahn nonlinearity. (k) Optimization convergence of the inverse reconstruction. (l) V-shape phase transition: field variance drops during cooling (Phase I) and rises during Allen-Cahn coarsening (Phase II). Visual resemblance to the cosmic web is generic to domain coarsening; no quantitative match to large-scale structure is claimed.

990 4.17. Combined verdict across three observational regimes

991 *Combined verdict: dense-scale unresolvability is a feature, not a falsification..*
 992 The naive reading of Pillars F, G, H, J is that they “rule out” DSC, since
 993 each individually gives χ^2/dof much larger than ΛCDM and $\Delta\chi^2$ values in

the hundreds-to-tens-of-thousands. We argue that this reading misses the central physical content of the $1/\ln^2$ ansatz, which is established by Pillars M (section 4.3) and N (section 4.2):

- (i) The $1/\ln^2(t/t_P)$ function has $8.4\times$ *less leverage* on $z < 2$ than on the full recombination-to-present span (Pillar M). Restricted to the dense-data regime, $\xi(z) = 10^5/\ln^2(t/t_P)$ varies by only $\Delta\xi \approx 0.10$, vs. $\Delta\xi \approx 0.84$ at anchor scale. In the dense regime DSC is therefore mathematically nearly degenerate with constant H_0 , regardless of the value of (H_∞, β) . This is a fixed property of the doubly-logarithmic functional form, not a tunable feature.
- (ii) Pillar N shows that the same property guarantees that the present-epoch drift is $|\dot{\alpha}/\alpha| = |\Gamma| \cdot 5.25 \times 10^{-22} \text{ yr}^{-1}$, four orders of magnitude below the Rosenband 2008 atomic-clock bound at the largest $|\Gamma|$ supported by Pillar A. The $1/\ln^2$ ansatz is therefore *automatically* consistent with present-epoch laboratory bounds, not by construction or fine-tuning, but as a generic property of the functional form. Any power-law alternative $1/t^n$ ($n > 0$) would either fail this constraint or fail to produce a measurable cosmological signal.

The dense-data “rejection” is therefore the predicted behavior: the model says $H(z)$ is nearly constant for $z < 2$, and the data confirm it. The model says $|\dot{\alpha}/\alpha|$ is unobservable in the laboratory, and the laboratory confirms it. The model says cosmologically-large signals are barely observable in H_0 anchors and quasar absorption, and Pillars A and C find such signals at $\sim 5\text{--}8\sigma$. All three regimes are consistent with the same $1/\ln^2$ ansatz.

What we genuinely retract is narrower: (a) the CMB lattice surrogate of sections 4.5 to 4.16 does *not* reproduce the Planck angular spectrum (section 4.14); this is a limitation of the forward implementation, not the $1/\ln^2$ ansatz. (b) The Pillar A Keck signal at $\Gamma \approx +6$, while statistically significant in two independent Keck datasets, is decomposed by Pillar K into an instrumental Keck–VLT systematic; the cosmological drift, if any, is bounded by $|\Gamma| \lesssim 1.5$ (Full Keck+VLT union under budget i) and is fully consistent with Pillar B (King 2012 full sample, budget ii: $\Gamma = +0.14 \pm 0.98$) and with the atomic-clock bound (Pillar N).

We therefore present the $1/\ln^2$ ansatz as a *self-consistent, internally-bounded, frame-spanning* parameterization of slow time-variation in cosmo-

logical parameters, supported by anchor-scale evidence (Pillars A, C), confirmed by the predicted dense-scale unresolvability (Pillars F, G, H, J), and bounded by laboratory atomic-clock measurements (Pillar N).

5. Discussion

What the Gaussianity finding does and does not show.. The near-Gaussian one-point statistics of the lattice are consistent with two simple expectations rather than a non-trivial emergence. First, the initial conditions are random Gaussian fields, and the Störmer–Verlet update is linear in ϕ in the parameter range we explore; a linear evolution of a Gaussian field is also Gaussian. Second, each pixel of the projected sky is a sum over many lattice modes, and the central limit theorem favors Gaussianity regardless of the individual mode statistics. The contrast with the dissipative CML baseline (kurtosis = -1.60) shows that strongly nonlinear dynamics destroy this Gaussianity, but the comparison cannot isolate symplecticity as the active ingredient: the two baselines differ in nonlinearity, conservation properties, and integration order simultaneously.

Lattice peaks vs. baryon–photon acoustic oscillations.. The peaks in the lattice $D(k)$ arise from standing-wave physics on a finite-volume lattice with a time-dependent effective propagation speed. On a sphere of radius R inside the lattice they project, in the flat-sky limit, to multipoles $\ell \approx 2\pi R_{\text{shell}} k$. With the geometry we use this gives $\ell \approx 30, 60, 90$ for the first three lattice peaks, far from the Planck multipoles $\ell \approx 220, 540, 810$. The peak positions in ℓ -space drift with the projection radius and the lattice size (section 4.14), which is the behavior expected from standing waves rather than from physical acoustic oscillations. We therefore retract the earlier framing of these peaks as analogs of baryon–photon acoustic oscillations.

A note on retroactive scale fitting (capability we deliberately do not exercise).. A common practice in lattice and effective-model cosmology is to introduce nuisance parameters such as the projection radius R_{shell} , the lattice spacing a_{lat} , the effective sound speed c_s^2 , or the Silk-damping scale k_s as free fit parameters when comparing to Planck C_ℓ . With three to four such nuisance parameters jointly fit by MCMC, the lattice peak positions of section 4.7 could be aligned with the Planck peaks at $\ell \approx 220, 540, 810$, and the corresponding χ^2/dof would fall from ~ 80 (the unfitted value of section 4.14) to below ~ 5 . Our differentiable JAX engine has this capability: a 30-minute

GPU MCMC over four nuisance parameters with the existing forward model would deliver such a fit. We deliberately do *not* do this, for two reasons. First, the lattice peaks shift their ℓ positions linearly with R_{shell} (we explicitly verified this in section 4.14 via the projection-radius and grid-size sensitivity scan); a successful retroactive fit therefore *confirms* the standing-wave interpretation rather than refuting it, since the fit succeeds precisely because the geometric standing-wave scale was tuned to match the Planck peaks. Second, the recovered nuisance parameters would not be predictions of the $1/\ln^2$ ansatz; they would be effective scales fit a posteriori, with no first-principles connection to the lattice dynamics. We report the unfitted, geometry-fixed $\chi^2/\text{dof} \approx 80$ as the honest test of whether the unmodified $1/\ln^2$ lattice forward model reproduces the Planck spectrum, and the answer is no. A reader interested in the post-fit χ^2/dof value can reproduce it from the existing notebooks; the engine can do it, we choose not to claim it.

Status of the connection to the DSC framework.. The simulations here implement the $1/\ln^2$ cooling law as a definite ansatz to study, not as a derived prediction of Planck-scale physics. The motivation from Riemann zero spectral statistics is treated here as a working motivation for the $1/\ln^2$ functional form of the cooling law, not as an established physical correspondence. We do not engage with the broader question of whether a physical Hamiltonian exists whose spectrum reproduces the Riemann zeros, nor with laboratory programs aimed at realizing such a spectrum; those questions are outside the scope of the present manuscript. Within the lattice surrogate we observe that the $1/\ln^2(t)$ cooling controls the location of the threshold-defined effective freeze-out scale (section 4.8), but this scale is sensitive to threshold and grid parameters and we do not interpret it as an analog of the comoving sound horizon.

Comparison with related approaches.. Coupled map lattices have been studied as models for spatiotemporal chaos and pattern formation [31], but not in a cosmological context with adiabatic cooling. Lattice cosmology approaches [40] discretize Einstein’s equations directly, while causal set theory [41] replaces the manifold with a partial order. The DSC lattice differs from both: it does not attempt to recover general relativity from the lattice, and it should not be read as an alternative formulation of cosmology. The differentiable simulation-based inference component connects to the growing field of likelihood-free inference in cosmology [22, 21], where forward

simulations replace analytical likelihoods, and to differentiable cosmological
simulators [18, 19, 20].

Comparison to existing time-varying H_0 analyses, and the anchor-vs-dense leverage asymmetry. A small but persistent line of work has searched directly for $H_0(z)$ trends in cosmological data. Dainotti *et al.* 2021 [42] binned the Pantheon Type Ia supernova compilation in redshift and reported a $\sim 2\sigma$ descending trend $H_0(z) = \tilde{H}_0/(1+z)^\alpha$ with $\alpha \approx 0.01$. Krishnan *et al.* 2020 [43] found a similar mild descending trend combining Pantheon, BAO, and cosmic-chronometer data ($\sim 2\sigma$). Kazantzidis & Perivolaropoulos 2020 [44] report $\sim 2\sigma$ on Cepheid+TRGB distance ladders. Colgáin *et al.* 2022 [45] extended the analysis with quasar+BAO data and again obtained $\lesssim 2\sigma$. The Hubble-tension model census of Schöneberg *et al.* 2022 [10] ranks dozens of solutions (early dark energy, modified recombination, etc.) and notes that none currently survive joint multi-probe scrutiny.

Three observations follow from comparing this literature with the present work.

(i) *Functional form.* Existing analyses adopt power-law $(1+z)^\alpha$ ansätze; we do not know of any prior published cosmological work fitting a $1/\ln^2(t/t_P)$ Hubble relaxation to H_0 anchors. The functional form here is independent of the literature ansätze, motivated by the Riemann-zeta-zero spectral density (section 2.1) and by the slowness requirement of section 4.2.

(ii) *Leverage choice.* The literature analyses operate exclusively within the $z < 2.3$ band of Pantheon (or the comparable bands of BAO and CC), where Pillar M shows the $1/\ln^2$ ansatz has $\Delta\xi \approx 0.10$ of dynamic range. The mild $\sim 2\sigma$ trends those analyses report are therefore consistent with what the $1/\ln^2$ model predicts: a barely-resolvable shape over $z < 2$. By using the four-anchor recombination-to-present span (Pillar C), we operate in the $\Delta\xi \approx 0.84$ regime where the same ansatz has $8.4\times$ the leverage—hence $R^2 = 0.91, \chi^2/\text{dof} = 0.28$ rather than $\sim 2\sigma$. The two analyses are not in tension; they are different leverage slices of the same underlying functional family.

(iii) *The anchor-vs-dense asymmetry is a generic property of the field.* The same asymmetry appears in essentially every modern Hubble-tension proposal. Early Dark Energy [10] resolves the SH0ES–Planck anchor difference by modifying recombination physics, but is in tension with low-redshift growth-rate measurements (σ_8 tension; the so-called yin-yang problem). CPL (w_0w_a) achieves a marginal $\Delta\text{BIC} = -2.6$ on dense BAO+SN+CC data

1137 (Pillar J) but does not address the anchor-level H_0 tension by construc-
 1138 tion. Modified-recombination scenarios match Planck but introduce tensions
 1139 elsewhere. The DSC $1/\ln^2$ ansatz, like all of these proposals, exhibits a
 1140 leverage-dependent fit quality. What is, to our knowledge, novel here is
 1141 making the asymmetry mathematically explicit (Pillar M, $8.4\times$ ratio fixed
 1142 by the doubly-logarithmic form), bounding it by laboratory atomic-clock
 1143 measurements (Pillar N), and providing a quantitative falsification roadmap
 1144 (Pillar O, $\gtrsim 11\sigma$ discrimination from ~ 5 JWST cosmic-chronometer points
 1145 at $z \in [2, 5]$). Whether this self-consistency framing extends to other modern
 1146 proposals (e.g. EDE) is, in our view, a worthwhile direction for the field.

1147 *What is forward prediction and what is inverse fit (a guide for the reader)..*
 1148 The CMB-related figures of this manuscript fall in two qualitatively dif-
 1149 ferent categories. *Forward predictions* (fig. A.31a in the appendix; fig. 20;
 1150 fig. 23) report the unmodified output of the DSC lattice: random Gaus-
 1151 sian initial conditions, evolved by weakly damped Störmer–Verlet steps with
 1152 $1/\ln^2$ cooling, projected onto a HEALPix sphere. No fit parameters enter
 1153 the forward simulation, except possibly an overall amplitude rescaling in the
 1154 spectral comparison of section 4.14. These are first-principles predictions of
 1155 the framework—they show what $1/\ln^2$ cosmology, taken at face value, gener-
 1156 ates as a sky map and as an angular spectrum. The honest verdict on these
 1157 forward predictions is the one we already report: pixel one-point distributions
 1158 are close to Planck SMICA (section 4.10), but the angular power spectrum
 1159 gives $\chi^2/\text{dof} \approx 80$ vs. CAMB Planck-2018 (section 4.14), and the lattice
 1160 $D(k)$ -peaks fall at $\ell \approx 30, 60, 90$ rather than the observed $\ell \approx 220, 540, 810$
 1161 (Pillar K). The first-handed prediction of DSC is therefore that the lattice
 1162 surrogate, taken alone, does *not* reproduce the observed CMB at the spec-
 1163 trum level.

1164 *Inverse reconstructions* (fig. 29 achieving in-sample $r = 0.98$) are some-
 1165 thing fundamentally different. There the differentiable engine optimizes
 1166 $\sim 10^6$ voxel parameters of the 3D primordial field to match $\sim 10^4$ ob-
 1167 served sphere-projected pixels. With $\sim 100\times$ more free parameters than data
 1168 points, reaching high in-sample correlation is mathematical inevitability, not
 1169 physical reconstruction. The held-out hemisphere test (fig. 26, $r = -0.17$) is
 1170 the appropriate diagnostic and we treat it as such throughout the manuscript.
 1171 The inverse figures should therefore not be read as “DSC reproduces the
 1172 Planck CMB” — they are demonstrations that the differentiable engine works
 1173 as advertised, not predictions of the framework.

1174 We separate these two categories explicitly because conflating them would
 1175 be the most likely misreading of this manuscript. The forward figures show
 1176 what $1/\ln^2$ *actually predicts*; the inverse figures show what the JAX engine
 1177 *can be made to fit*. Both are useful: the former delineates the empirical reach
 1178 (and limit) of the simple ansatz, while the latter establishes the engine as a
 1179 methodological building block for simulation-based inference workflows that
 1180 pair lattice forward models with multi-probe data and explicit volumetric
 1181 priors.

1182 *Forward identifiability vs. field-level underdetermination..* The forward twin
 1183 experiment (section 4.9) optimizes 3 global physics parameters against ~ 75
 1184 power-spectrum bins, an overdetermined system that recovers the param-
 1185 eters up to a c^2 - k_{opt} degeneracy direction. The voxel-level reconstruction
 1186 (section 4.15), in contrast, attempts to recover $\sim 10^6$ free parameters from
 1187 $\sim 10^4$ sky pixels. The high in-sample correlation $r = 0.98$ is not surpris-
 1188 ing at this ratio of unknowns to data, and the held-out hemisphere result
 1189 $r = -0.17$ is the appropriate diagnostic: the inverse problem is severely un-
 1190 derdetermined and the recovered 3D field cannot be interpreted as a physical
 1191 primordial reconstruction. We therefore present the differentiable engine—
 1192 gradients flowing cleanly through 25 Störmer–Verlet steps and the sphere-
 1193 projection layer—as a methodological building block for future multi-probe
 1194 inference workflows that pair lattice forward models with volumetric priors,
 1195 rather than as a tool for inferring the 3D primordial state from CMB data
 1196 alone.

1197 *The continuation past the CMB-epoch step is a phenomenological visualiza-*
 1198 *tion..* The V-shape transition in field variance (section 4.16) is produced
 1199 by a $1/\ln^2(t)$ cooling phase followed by an Allen–Cahn coarsening phase.
 1200 The Allen–Cahn nonlinearity is a generic toy potential, not a gravitational
 1201 interaction, and the resulting domain morphology is generic to coarsening
 1202 dynamics. We retract the earlier suggestion of an analogy with cosmic infla-
 1203 tion: the smoothing in our run is driven by an effective sound-speed decay
 1204 rather than by a slow-roll scalar field, and the analogy is not quantitative.
 1205 The cosmic-timeline figure should be read as an illustration of the lattice’s
 1206 behavior under a chosen toy nonlinearity, not as a model of cosmic structure
 1207 formation.

6. Limitations

Several limitations must be acknowledged.

1. **Kurtosis tension.** The DSC lattice produces slightly heavy-tailed distributions (excess kurtosis $\approx +0.3$), while the Planck SMICA map has light tails (≈ -0.5). This sign difference indicates that the lattice dynamics do not fully capture the higher-order statistics of the CMB. The weak nonlinear term ($-\alpha\phi^2$) used in 3D simulations contributes to this discrepancy; reducing α improves the kurtosis but at the cost of less pronounced cosmic web structure.
2. **Three free parameters.** Matching the acoustic peak positions to Λ CDM requires fitting c_{base}^2 , η , and k_{opt} . The DSC framework does not predict their values from first principles (beyond the conjecture $k_{\text{opt}} \approx \alpha_F/2$). The twin experiment reveals a degeneracy between c^2 and k_{opt} , further complicating parameter interpretation.
3. **Mock Λ CDM comparison.** Our power spectrum comparisons use the analytic mock eq. (16) rather than the full Planck best-fit C_ℓ . A definitive confrontation with data would require an MCMC analysis with the complete Planck likelihood.
4. **Surrogate lattice.** The 2D and 3D simulations are effective models, not direct implementations of the Planck-scale lattice $\mathbb{Z}^3 \times \mathbb{N}$. The connection between lattice-unit quantities (e.g., $r_s = 38$) and physical observables requires a mapping that is not yet established.
5. **No polarization or tensor modes.** We study only scalar (temperature) fluctuations. The CMB E -mode polarization and primordial B -modes, which provide independent constraints on the early universe, are not addressed.
6. **Conjectural physical basis.** The derivation chain from the assumed $1/\ln$ spectral density to the $1/\ln^2$ cooling law involves three conjectural steps (Assumptions 1–3 in section 2.1). These are physically motivated but not rigorously proven. The present work should be viewed as exploring the *consequences* of these assumptions in a computational setting, not as establishing their validity.

- 1240 **7. Inverse reconstruction does not generalize.** The held-out test
 1241 (section 4.15) shows that the 3D reconstruction from a single hemi-
 1242 sphere does not predict the unseen hemisphere ($r_{\text{test}} = -0.17$). The
 1243 reconstruction demonstrates differentiable-engine capability, not phys-
 1244 ical uniqueness of the inferred 3D structure.
- 1245 **8. Quantitative C_ℓ mismatch.** The forward simulation yields $\chi^2/\text{dof} \approx$
 1246 3.3 against Planck ($\ell = 2\text{--}50$), far from a statistically acceptable fit.
 1247 The acoustic peak structure visible in the 2D $D(k)$ is partly lost upon
 1248 sphere projection, and a full Planck likelihood (MCMC) analysis has
 1249 not been performed.

1250 7. Conclusion

1251 We presented a consolidated empirical and computational evaluation of
 1252 the Discrete Symplectic Cosmology (DSC) framework, in which a $1/\ln^2(t)$
 1253 adiabatic cooling law—inspired by the spectral statistics of the Riemann
 1254 zeta zeros—is the central ansatz tying microscopic non-autonomous dynamics
 1255 to time-dependent cosmological parameters. Across thirteen tests at three
 1256 observational scales, the verdict is that the $1/\ln^2$ ansatz is *self-consistent*
 1257 *across regimes*.

1258 *The $1/\ln^2(t)$ ansatz is the unique slow-drift functional form that reconciles*
 1259 *three observational regimes..* Any time variation of fundamental constants
 1260 must satisfy three constraints simultaneously: (i) be observable at cosmologically-
 1261 large time spans (where Pillars A and C find $> 5\sigma$ signals); (ii) be unresolv-
 1262 able from constant on dense $z < 2$ data (where Pillars F, G, H, J return
 1263 $\chi^2/\text{dof} \gg 1$ for any non-constant ansatz); (iii) be undetectable in present-
 1264 epoch laboratory atomic clocks (Rosenband 2008 bound). Pillar M quantifies
 1265 the $8.4\times$ leverage difference between regimes (i) and (ii), and Pillar N quan-
 1266 tifies the four orders of magnitude of headroom relative to (iii). A power-law
 1267 $1/t^n$ with $n > 0$ would fail at least one of these three constraints. The doubly-
 1268 logarithmic $1/\ln^2(t/t_P)$ ansatz uniquely meets all three, providing physical
 1269 motivation for the central functional form independent of the Riemann-zeta
 1270 motivation.

1271 *Anchor-scale evidence..* Pillar C: the 4-anchor H_0 relaxation across SH0ES,
 1272 JWST, Planck, and BAO yields $R^2 = 0.91$, $\chi^2/\text{dof} = 0.28$, recasting the
 1273 SH0ES–Planck Hubble tension as a smooth time-dependent relaxation rather

1274 than a discrepancy. Pillar A: on the Murphy 2003 128-absorber Keck/HIRES
1275 catalog, the $1/\ln^2(t)$ drift is detected at $\Gamma = +5.56 \pm 0.99$ (5.6σ , $\Delta\chi^2 = 31.83$
1276 under per-fit-error-only weighting; 2.81σ under Murphy's σ_{rand} -augmented
1277 total budget). Pillar K cross-validates this on a second independent Keck/HIRES
1278 dataset (King 2012 Keck-only, $\Gamma = +6.28 \pm 0.97$, 6.5σ); the joint Murphy-03
1279 + King-12 Keck-only union reaches $\Gamma = +5.92 \pm 0.69$ ($+8.6\sigma$, $\Delta\chi^2 = 73.5$
1280 under budget i).

1281 *Dense-scale evidence (predicted unresolvability)..* On 32 cosmic chronometers
1282 (Pillar F), Pantheon+ (1580 SNe; Pillar G), DESI DR2 BAO (Pillar H), and
1283 the joint 1623-constraint dataset (Pillar J), DSC is degenerate with constant
1284 H_0 over $z < 2$, returning χ^2/dof values in the tens. This is the predicted
1285 behavior, not a falsification: Pillar M shows that the $1/\ln^2(t)$ ansatz has only
1286 $\Delta\xi \times 10^5 = 0.10$ leverage on $z < 2$ vs. 0.84 on the recombination-to-now span,
1287 a fixed ratio determined entirely by the doubly-logarithmic functional form.
1288 CPL ($w_0 w_a$) is mildly preferred over ΛCDM at $\Delta\text{BIC} = -2.6$, providing
1289 independent evidence that the late-time universe departs from ΛCDM , but
1290 DSC by construction does not compete with $w_0 w_a$ as a description of *dense-*
1291 *scale* departure.

1292 *Instrumental decomposition (Pillar K)..* The King 2012 VLT-only subset
1293 gives Γ of opposite sign at -3.7σ (budget i) or -1.8σ (budget ii); adding
1294 VLT to the Keck-only union reduces the apparent significance to $+3.0\sigma$ (i)
1295 or $+1.5\sigma$ (ii). We agree with the wavelength-calibration findings of Whitmore
1296 & Murphy 2015 [34] and Wilczynska *et al.* 2020 [35]: the apparent Pillar A/K
1297 Keck-only signal at $\Gamma \approx +6$ is driven by Keck instrumental systematics; the
1298 cosmologically-allowed drift, after the VLT subset is included, is bounded by
1299 $|\Gamma| \lesssim 1.5$ (still well within the atomic-clock bound of Pillar N).

1300 *What we retract..* We retract: (i) any quantitative claim that the lattice for-
1301 ward model reproduces the Planck C_ℓ at the level of individual multipoles
1302 ($\chi^2/\text{dof} \approx 80$; this is a limitation of the lattice surrogate, not the $1/\ln^2$
1303 ansatz), and (ii) any structural-formation interpretation of the V-shape lat-
1304 tice continuation past the CMB-epoch step (the Allen–Cahn nonlinearity is
1305 not gravity).

1306 *What survives, what is contributed..* Five contributions of this work survive
1307 the more conservative reading: (i) the $1/\ln^2(t/t_P)$ ansatz, motivated in-
1308 dependently by the Riemann-zeta-zero spectral density and by the require-
1309 ment of slowness across atomic-clock + cosmological + dense-data scales; (ii)

1310 anchor-scale empirical detection of the predicted drift in the Hubble param-
1311 eter (Pillar C, $R^2 = 0.91$) and in the fine-structure constant (Pillars A, K,
1312 $+8.6\sigma$ in the Keck-only union); (iii) explicit demonstration of dense-scale un-
1313 resolvability (Pillars F, G, H, J), with the leverage diagnostic (Pillar M) and
1314 atomic-clock consistency (Pillar N) showing that this is the predicted behav-
1315 ior; (iv) the instrument-resolved decomposition of the apparent quasar signal
1316 (Pillar K), establishing that Keck-only datasets give $\Gamma \approx +6$ but VLT-only
1317 gives Γ of opposite sign — a textbook signature of instrumental systematics;
1318 (v) the differentiable JAX lattice surrogate, with clean gradients through 25
1319 Störmer–Verlet steps and a sphere-projection layer, useful as a methodolog-
1320 ical building block for simulation-based inference workflows.

1321 *Future directions.* The most informative next steps follow directly from the
1322 Pillar O forecast: (i) targeting JWST high- z passive galaxies [28] for cosmic-
1323 chronometer measurements at $z \in [2, 10]$, which is the regime where the DSC
1324 and Λ CDM Hubble curves diverge by 65–95% and as few as five points at 10%
1325 precision suffice for $\gtrsim 11\sigma$ discrimination; (ii) replacing the Pillar A Keck
1326 datasets with VLT/UVES re-observations of the same lines of sight to test
1327 whether the cosmologically-allowed $|\Gamma| < 1.5$ band can be tightened; (iii)
1328 using ELT-class instruments (HARPS-3, ANDES) to push the atomic-line
1329 laboratory bound on $|\dot{\alpha}/\alpha|$ from 10^{-17} yr^{-1} towards 10^{-19} yr^{-1} , which would
1330 tighten the Pillar N bound on $|\Gamma|$ by two orders of magnitude; (iv) replacing
1331 the compressed CMB observables of Pillar D with a full Planck-likelihood
1332 analysis sensitive to the recombination-to-present anchor leverage; (v) gen-
1333 eralizations of the lattice forward model that may improve the Pillar 10 C_ℓ
1334 comparison without modifying the underlying $1/\ln^2$ cooling law.

1335 The doubly-logarithmic functional form is, in our view, the right family
1336 to develop further; the lattice forward implementation is the part that needs
1337 work; and the JWST high- z Hubble program is the experiment that decides
1338 whether the DSC ansatz is the candidate solution to the modern dark clouds,
1339 or merely a useful illustration of how a slow-drift framework should be tested.

1340 Acknowledgments

1341 The Planck SMICA map used in this work is based on observations ob-
1342 tained with Planck (<http://www.esa.int/Planck>), an ESA science mission
1343 with instruments and contributions directly funded by ESA Member States,
1344 NASA, and Canada.

Data Availability Statement

Code associated with this project is available at https://github.com/maris205/riemann_engine. The Planck SMICA map (COM_CMB_IQU-smica_2048_R3.00_full.fits) is publicly available from the ESA Planck Legacy Archive.

Appendix A. Overview Figure

Figure A.31 provides an overview of the main results.

References

- [1] Planck Collaboration, Planck 2018 results. I. Overview and the cosmological legacy of Planck, *Astronomy & Astrophysics* 641 (2020) A1. doi:10.1051/0004-6361/201833880.
- [2] Planck Collaboration, Planck 2018 results. VI. Cosmological parameters, *Astronomy & Astrophysics* 641 (2020) A6. doi:10.1051/0004-6361/201833910.
- [3] Planck Collaboration, Planck 2018 results. IX. Constraints on primordial non-Gaussianity, *Astronomy & Astrophysics* 641 (2020) A9. doi:10.1051/0004-6361/201935891.
- [4] Planck Collaboration, Planck 2018 results. V. CMB power spectra and likelihoods, *Astronomy & Astrophysics* 641 (2020) A5. doi:10.1051/0004-6361/201936386.
- [5] S. Dodelson, F. Schmidt, *Modern Cosmology*, 2nd Edition, Academic Press, 2020.
- [6] Planck Collaboration, Planck 2018 results. X. Constraints on inflation, *Astronomy & Astrophysics* 641 (2020) A10. doi:10.1051/0004-6361/201833887.
- [7] V. Mukhanov, *Physical Foundations of Cosmology*, Cambridge University Press, 2005. doi:10.1017/CB09780511790553.
- [8] R. Durrer, *The Cosmic Microwave Background*, Cambridge University Press, 2008. doi:10.1017/CB09780511817205.




Figure A.31: Overview of DSC weakly damped symplectic lattice surrogate. (a) Full-sky temperature fluctuation map from 3D Störmer–Verlet evolution (128^3 lattice, 60 steps) projected onto a Mollweide sphere by sampling a radius- $0.35 N_{\text{grid}}$ shell. The visible large-scale red–blue patches are lattice standing-wave modes projected onto the sphere (the lattice has a finite size, and the doubly-logarithmic cooling damps small-scale modes preferentially); they are *not* physical CMB cold spots, and the lattice $D(k)$ peaks do not map to the Planck multipoles (section 4.14, Pillar K). (b) Scaled power spectrum $D(k) = k^2 P(k)$ with multimodal oscillation peaks (blue triangles), compared to a toy analytic reference (dashed); the toy reference is used only for shape-level comparison and is replaced by the real CAMB Planck-2018 D_ℓ^{TT} benchmark in section 4.14. (c) Ablation: pixel distribution of dissipative CML (red) vs. weakly damped symplectic lattice (blue); only the symplectic-family lattice matches a Gaussian reference (black dashed), but as discussed in section 4.5 this comparison cannot isolate symplecticity as the unique cause. (d) Pixel one-point distribution of the DSC 3D forward simulation (blue) is close to that of the actual Planck SMICA map (purple), but at the angular-spectrum level the forward model gives $\chi^2/\text{dof} \approx 80$ vs. the CAMB Planck-2018 C_ℓ (section 4.14); the visual resemblance at the pixel level should not be over-interpreted as a CMB-physics match.

- [9] E. Di Valentino, O. Mena, S. Pan, L. Visinelli, W. Yang, A. Melchiorri, D. F. Mota, A. G. Riess, J. Silk, In the realm of the Hubble tension—a review of solutions, *Classical and Quantum Gravity* 38 (2021) 153001. doi:10.1088/1361-6382/ac086d.
- [10] N. Schöneberg, G. Franco Abellán, A. Pérez Sánchez, S. J. Witte, V. Poulin, J. Lesgourgues, The H0 Olympics: A fair ranking of proposed models, *Physics Reports* 984 (2022) 1–55. doi:10.1016/j.physrep.2022.07.001.
- [11] DESI Collaboration, DESI 2024 VI: Cosmological Constraints from the Measurements of Baryon Acoustic Oscillations, arXiv preprint (2024). arXiv:2404.03002.
- [12] J.-P. Uzan, The fundamental constants and their variation: observational and theoretical status, *Reviews of Modern Physics* 75 (2003) 403–455. doi:10.1103/RevModPhys.75.403.
- [13] C. J. A. P. Martins, The status of varying constants: a review of the physics, searches and implications, *Reports on Progress in Physics* 80 (2017) 126902. doi:10.1088/1361-6633/aa860e.
- [14] L. Wang, The emergence of prime distribution from low-dimensional deterministic chaos, *Research in Mathematics* 13 (1) (2026). doi:10.1080/27684830.2026.2684334.
- [15] L. Wang, Spectral Isomorphism between Renormalization Flow in Non-Autonomous Quadratic Maps and Riemann Zeros, v4.0 (2026). doi:10.5281/zenodo.19045440.
- [16] E. Hairer, C. Lubich, G. Wanner, *Geometric Numerical Integration: Structure-Preserving Algorithms for Ordinary Differential Equations*, 2nd Edition, Springer, 2006. doi:10.1007/3-540-30666-8.
- [17] B. Leimkuhler, S. Reich, *Simulating Hamiltonian Dynamics*, Cambridge University Press, 2004. doi:10.1017/CB09780511614118.
- [18] C. Modi, F. Lanusse, U. Seljak, FlowPM: Distributed TensorFlow implementation of the FastPM cosmological N-body solver, *Astronomy and Computing* 37 (2021) 100505. doi:10.1016/j.ascom.2021.100505.

- [19] Y. Li, L. Lu, U. Seljak, Differentiable cosmological simulation with the adjoint method, *The Astrophysical Journal Supplement Series* 259 (2022) 36. doi:10.3847/1538-4365/ac5f4a.
- [20] J.-E. Campagne, F. Lanusse, J. Zuntz, et al., JAX-COSMO: An End-to-End Differentiable and GPU Accelerated Cosmology Library, *The Open Journal of Astrophysics* 6 (2023) 15. doi:10.21105/astro.2302.05163.
- [21] K. Cranmer, J. Brehmer, G. Louppe, The frontier of simulation-based inference, *Proceedings of the National Academy of Sciences* 117 (2020) 30055–30062. doi:10.1073/pnas.1912789117.
- [22] J. Alsing, B. Wandelt, S. Feeney, Massive optimal data compression and density estimation for scalable, likelihood-free inference in cosmology, *Monthly Notices of the Royal Astronomical Society* 488 (2019) 4440–4458. doi:10.1093/mnras/stz1960.
- [23] D. G. Figueroa, A. Florio, F. Torrenti, W. Valkenburg, The art of simulating the early Universe. Part I. Integration techniques and canonical cases, *Journal of Cosmology and Astroparticle Physics* 2021 (2021) 035. doi:10.1088/1475-7516/2021/04/035.
- [24] A. Caravano, E. Komatsu, A. Linde, K. D. Lozanov, Lattice simulations of inflation, *arXiv preprint* (2022). arXiv:2209.13616.
- [25] T. Rosenband, D. B. Hume, P. O. Schmidt, C. W. Chou, A. Brusch, L. Lorini, W. H. Oskay, R. E. Drullinger, T. M. Fortier, J. E. Stalnaker, S. A. Diddams, W. C. Swann, N. R. Newbury, W. M. Itano, D. J. Wineland, J. C. Bergquist, Frequency ratio of Al^+ and Hg^+ single-ion optical clocks; metrology at the 17th decimal place, *Science* 319 (2008) 1808–1812. doi:10.1126/science.1154622.
- [26] A. G. Riess, et al., A Comprehensive Measurement of the Local Value of the Hubble Constant with $1 \text{ km s}^{-1} \text{ Mpc}^{-1}$ Uncertainty from the Hubble Space Telescope and the SH0ES Team, *The Astrophysical Journal Letters* 934 (2022) L7. doi:10.3847/2041-8213/ac5c5b.
- [27] A. G. Riess, et al., JWST Observations Reject Unrecognized Crowding of Cepheid Photometry as an Explanation for the Hubble Tension at 8σ Confidence, *The Astrophysical Journal Letters* 962 (2024) L17. doi:10.3847/2041-8213/ad1ddd.

- 1437 [28] K. Glazebrook, T. Nanayakkara, C. Schreiber, C. Lagos, L. Kawin-
1438 wanichakij, C. Jacobs, H. Chittenden, G. Brammer, G. G. Kacprzak,
1439 I. Labbé, D. Marchesini, Z. C. Marsan, C. Papovich, R. F. Quadri,
1440 R.-S. Remus, K.-V. Tran, J. Esdaile, B. Forrest, A massive galaxy
1441 that formed its stars at $z \approx 11$, *Nature* 628 (2024) 277–281. doi:
1442 10.1038/s41586-024-07191-9.
- 1443 [29] P. E. Kloeden, M. Rasmussen, *Nonautonomous Dynamical Systems*,
1444 American Mathematical Society, 2011. doi:10.1090/surv/176.
- 1445 [30] K. M. Górski, E. Hivon, A. J. Banday, B. D. Wandelt, F. K. Hansen,
1446 M. Reinecke, M. Bartelmann, HEALPix: A Framework for High-
1447 Resolution Discretization and Fast Analysis of Data Distributed on
1448 the Sphere, *The Astrophysical Journal* 622 (2005) 759–771. doi:
1449 10.1086/427976.
- 1450 [31] K. Kaneko, Overview of coupled map lattices, *Chaos* 2 (1992) 279–282.
1451 doi:10.1063/1.165869.
- 1452 [32] T. Akiba, S. Sano, T. Yanase, T. Ohta, M. Koyama, Optuna: A Next-
1453 generation Hyperparameter Optimization Framework, in: *Proceedings*
1454 *of the 25th ACM SIGKDD International Conference on Knowledge Dis-*
1455 *covery & Data Mining*, 2019, pp. 2623–2631. doi:10.1145/3292500.
1456 3330701.
- 1457 [33] M. T. Murphy, J. K. Webb, V. V. Flambaum, Further evidence for a vari-
1458 able fine-structure constant from Keck/HIRES QSO absorption spectra,
1459 *Monthly Notices of the Royal Astronomical Society* 345 (2003) 609–
1460 638. arXiv:astro-ph/0306483, doi:10.1046/j.1365-8711.2003.
1461 06970.x.
- 1462 [34] J. B. Whitmore, M. T. Murphy, Impact of instrumental systematic errors
1463 on fine-structure constant measurements with quasar spectra, *Monthly*
1464 *Notices of the Royal Astronomical Society* 447 (2015) 446–462. doi:
1465 10.1093/mnras/stu2420.
- 1466 [35] M. R. Wilczynska, J. K. Webb, M. Bainbridge, J. D. Barrow, S. E. I.
1467 Bosman, R. F. Carswell, M. P. Dabrowski, V. Dumont, C.-C. Lee,
1468 A. C. Leite, K. Leszczyńska, J. Liske, K. Marosek, C. J. A. P. Mar-
1469 tins, D. Milaković, P. Molaro, L. Pasquini, Four direct measurements

1470 of the fine-structure constant 13 billion years ago, *Science Advances* 6
1471 (2020) eaay9672. doi:10.1126/sciadv.aay9672.

1472 [36] M. Moresco, L. Amati, L. Amendola, S. Birrer, J. P. Blakeslee,
1473 M. Cantiello, A. Cimatti, J. Darling, M. Della Valle, M. Fishbach,
1474 C. Grillo, N. Hamaus, D. Holz, L. Izzo, H. Jia, A. Kim, L. Lom-
1475 briser, A. Pourtsidou, I. Tutusaus, L. Verde, Unveiling the Universe with
1476 emerging cosmological probes, *Living Reviews in Relativity* 25 (2022) 6.
1477 doi:10.1007/s41114-022-00040-z.

1478 [37] D. Scolnic, D. Brout, A. Carr, A. G. Riess, et al., The Pantheon+
1479 Analysis: The Full Data Set and Light-curve Release, *The Astrophysical*
1480 *Journal* 938 (2022) 113. doi:10.3847/1538-4357/ac8b7a.

1481 [38] D. Brout, D. Scolnic, B. Popovic, A. G. Riess, et al., The Pantheon+
1482 Analysis: Cosmological Constraints, *The Astrophysical Journal* 938
1483 (2022) 110. doi:10.3847/1538-4357/ac8e04.

1484 [39] DESI Collaboration, DESI DR2 Results II: Measurements of Baryon
1485 Acoustic Oscillations and Cosmological Constraints, *arXiv preprint*
1486 (2025). arXiv:2503.14738.

1487 [40] T. Regge, General relativity without coordinates, *Il Nuovo Cimento* 19
1488 (1961) 558–571. doi:10.1007/BF02733251.

1489 [41] L. Bombelli, J. Lee, D. Meyer, R. D. Sorkin, Space-time as a causal
1490 set, *Physical Review Letters* 59 (1987) 521–524. doi:10.1103/
1491 *PhysRevLett*.59.521.

1492 [42] M. G. Dainotti, B. De Simone, T. Schiavone, G. Montani, E. Ri-
1493 naldi, G. Lambiase, On the Hubble Constant Tension in the SNe Ia
1494 Pantheon Sample, *The Astrophysical Journal* 912 (2021) 150. doi:
1495 10.3847/1538-4357/abeb73.

1496 [43] C. Krishnan, E. Ó Colgáin, M. M. Sheikh-Jabbari, T. Yang, Is there
1497 an early Universe solution to Hubble tension?, *Physical Review D* 102
1498 (2020) 103525. doi:10.1103/PhysRevD.102.103525.

1499 [44] L. Kazantzidis, L. Perivolaropoulos, Hints of a local matter underdensity
1500 or modified gravity in the low z Pantheon data, *Physical Review D* 102
1501 (2020) 023520. doi:10.1103/PhysRevD.102.023520.

- 1502 [45] E. Ó Colgáin, M. M. Sheikh-Jabbari, R. Solomon, G. Bargiacchi,
1503 S. Capozziello, M. G. Dainotti, D. Stojkovic, Revealing intrinsic flat
1504 Λ CDM biases with standardizable candles, Physical Review D 106
1505 (2022) L041301. doi:10.1103/PhysRevD.106.L041301.